

Dream-RSI: Recursive Self-Improvement through Evolving Worlds

Tong Zheng^{1,2}, Xidong Wu¹, Zheng Zhang¹, Zhankui He³, Chaoyi Zhang¹, Benjamin Coleman³, Ruoqiao Wei¹, Di Bai³, Haolin Liu⁴, Rui Liu², Xue Wang¹, Yue Zhuan¹, Wang-Cheng Kang³, Renkai Xiang¹, Heng Huang², Xinwu Cheng¹ and Yunsong Guo¹

¹Google, ²University of Maryland, College Park, ³Google Deepmind, ⁴University of Virginia

Recursive self-improvement is becoming increasingly vital for autonomous AI agents, where progress hinges on discovering high-value solutions across complex domains. The driver of this process is effective exploration, however, managing and improving exploration strategies remains a major bottleneck. Current systems face a fundamental dilemma: fixed strategies fail to adapt as search spaces scale, while online policy optimization requires navigating vast meta-search spaces under delayed and expensive feedback over long-horizon rollouts. We introduce DREAM-RSI, a framework for scalable and recursively self-improving exploration. A lightweight orchestration layer makes exploration explicit and programmable while leaving the underlying coding agent unchanged. Our key insight is that accumulated discovery history can serve as a replay simulator over the realized search space. By performing dreaming in the replay simulator constructed from historical discovery trees, DREAM-RSI secures immediate, low-cost off-policy feedback to evaluate and refine exploration policies without invoking repetitive, expensive online evaluations. The improved policy is subsequently redeployed online to drive further discovery, continuously expanding the simulator pool in a self-improving loop. Across algorithm engineering, mathematical optimization, and GPU kernel engineering, DREAM-RSI achieves competitive or improved discovery quality while substantially reducing discovery cost in several settings.

 github.com/zhengkid/Dream-RSI |  dream-rsi.com

1. Introduction

Recursive self-improvement (RSI) has emerged as an ambitious goal for autonomous AI systems (Liu et al., 2026c). A common mechanism underlying RSI is an iterative discovery loop wherein agents generate candidate solutions, evaluate outcomes, incorporate feedback, and refine future iterations. Such discovery loops have driven substantial progress across scientific and algorithmic domains, including algorithm design (Novikov et al., 2025; Romera-Paredes et al., 2024), open-ended mathematical optimization (Anthropic, 2026; Georgiev et al., 2025), systems design (Cao et al., 2026; Jaber and Jaber, 2026), and agent self-improvement (Lee et al., 2026; Zhang et al., 2026b,c; Zheng et al., 2026a), with these discoveries increasingly feeding into the development of more capable AI systems. As agent capabilities improve and self-improvement targets become challenging, discovery increasingly requires long-horizon exploration over vast search spaces, often spanning thousands of proposal-evaluation cycles (OpenAI, 2026; Ye et al., 2026). At this scale, the ability to orchestrate exploration becomes critical (Zheng et al., 2026b). Poor exploration can waste substantial computation and time, severely limiting the efficiency and scalability of RSI.

Existing approaches have largely relied on manually designed exploration strategies that remain largely fixed throughout discovery (Du et al., 2026; Jiang et al., 2026; Novikov et al., 2025; Yan et al., 2026b; Ye et al., 2026). Fixed strategies cannot improve from accumulated discovery experience and may repeatedly allocate computation to ineffective search directions. Recent work therefore seeks to optimize exploration policies online during discovery (Liu et al., 2026a), but doing so faces




Figure 1 | **Overview of DREAM-RSI.** The system operates in a recursive self-improvement loop via three core stages: ① **Online Explore**, where the current exploration policy guides a coding agent to expand a discovery tree and log historical traces; ② **Construct Replay Simulator**, where the generated discovery tree is converted into a reusable simulator pool; and ③ **Dreaming-based Policy Improvement**, where the agent "dreams" up a massive pool of alternative policies in its mind. It then feeds these candidate policies into the replay simulator to simulate executions and derive rapid feedback, continuously refining its strategy (detailed in the Zoom-in box). The updated policy then redeplays for the next round of online exploration.

two fundamental bottlenecks. First, feedback is delayed and expensive at the meta level: unlike evaluating an individual candidate, assessing an exploration policy requires observing how it shapes the subsequent discovery process over many proposal–evaluation cycles. Second, the meta-policy space is vast: a newly proposed policy may perform poorly, so many alternatives may need to be tried. Together, these challenges make meta-level improvement particularly costly: each policy may require a long online rollout before receiving useful feedback, making it difficult to efficiently close the self-improvement loop at the exploration layer.

To address these bottlenecks, our key intuition is simple: a fast and inexpensive simulator of discovery would allow many exploration policies to be evaluated before costly online deployment. **Surprisingly, completed discovery histories already provide such a simulator.** While prior work treats past discovery history merely as static textual context (Hu et al., 2025; Ouyang et al., 2026b) or training data for weight fine-tuning (Wang et al., 2025; Yuksekogonul et al., 2026), a completed discovery process inherently records a structured tree of past exploration decisions and their realized code-execution outcomes. Drawing an analogy to model-based reinforcement learning and World Models (Ha and Schmidhuber, 2018; Hafner et al., 2023) (§2), once organized into a discovery tree, this history can serve as a **replay simulator**¹. As illustrated in Figure 2, an alternative exploration strategy can navigate this pre-recorded tree to traverse different subsets of recorded branches, in different orders, with different parallel groupings and stopping decisions. Because all execution outcomes are already saved in the tree, evaluating a new strategy requires only reading past records without rerunning the underlying discovery agent or evaluator. This transforms meta-policy improvement from an expensive online trial-and-error process into a fast, simulation-based “dreaming”

¹We use the terms replay simulator and worlds interchangeably.

procedure.

Building on this insight, we introduce DREAM-RSI, a framework for scalable and recursively self-improving meta-exploration in agent-driven discovery. We first make exploration explicit and programmable through a lightweight orchestration layer that controls branching, parallel exploration, and stopping while leaving the underlying coding agent unchanged. Rather than keeping this policy fixed, DREAM-RSI establishes a closed-loop self-improvement mechanism across three core stages (Figure 1): **(1) Online Exploration**, where the current policy guides real-world discovery and logs historical execution traces; **(2) Simulator Construction**, where recorded discovery trees are converted into a reusable replay simulator pool; and **(3) Dreaming-based Policy Improvement**, where candidate policies are evaluated via low-cost "dreaming" over the simulator. The updated policy is then redeployed online to generate new discovery experience and expand the simulator pool, closing a RSI loop at the meta-exploration layer.

Empirically, we evaluate DREAM-RSI across 8 scientific discovery tasks spanning three distinct domains: algorithm engineering, mathematical optimization, and GPU kernel engineering. In algorithm engineering (Lasso path solver), DREAM-RSI outperforms standard libraries like `sklearn` and strong baselines while reducing agent calls by up to $162\times$ over SIMPLETES and $1.7\times$ over fixed-exploration baselines. In mathematical optimization (sum-difference, autocorrelation, circle packing), it matches or surpasses strong baselines within 1k generations, yielding over $50\times$ budget savings compared to SIMPLETES. In GPU kernel engineering (KernelBench), it either reaches target execution speeds using $1.79\times$ – $2.43\times$ fewer generations or improves kernel performance by up to $2.09\times$ under identical budget constraints.

In summary, our main contributions are as follows: 1) **History as Replay Simulator**: We conceptualize completed discovery histories as replay simulators. This makes delayed exploration feedback reusable for efficient meta-exploration policy evaluation.; 2) **Meta-Layer RSI Loop (DREAM-RSI)**: We introduce DREAM-RSI, establishing a recursive self-improvement loop that continuously collects discovery histories through online exploration, constructs replay simulators from history to refine meta-exploration strategies via dreaming, and redeploys the upgraded policy online; 3) **Empirical Validation**: We conduct experiments to demonstrate that DREAM-RSI improves both discovery effectiveness and efficiency in several settings.

2. Motivation: Discovery History as a Replay Simulator

Consider an agent navigating toward a goal in an unfamiliar environment. During its first traversal, the agent may follow inefficient routes, encounter dead ends, backtrack, and gradually construct a map of the surrounding space. Once recorded, however, this experience becomes reusable: the resulting map supports planning without requiring the agent to physically revisit every location. A new navigation policy can instead reason over the accumulated map, avoid known dead ends, reconsider earlier decisions, and compare alternative routes before acting (Gupta et al., 2017).

This idea parallels model-based reinforcement learning (M. Moerland et al., 2023; Sutton, 1990). A model captures how an environment evolves in response to an agent’s actions, allowing policies to be trained or evaluated through simulated experience rather than repeated interaction with the real environment (Ha and Schmidhuber, 2018). The Dreamer family (Hafner et al., 2019, 2020, 2023, 2025) demonstrates this principle particularly clearly: an agent learns a compact dynamics model from collected experience and improves its policy by imagining trajectories within that model.

Long-horizon discovery admits an analogous structure. An exploration policy decides which directions to pursue, which candidates to refine, which branches to explore in parallel, and when




Figure 2 | **Discovery history as a replay simulator.** A deployed policy first explores online to generate a structured discovery tree containing historical execution traces (each node denote an attempt with its full observation). Thousands of candidate policies can then be tested within this simulator—evaluating alternative choices of search branches, exploration orders, concurrency levels, and stopping rules. **Since all node outcomes are pre-stored, a single costly online run enables thousands of rapid, zero-execution-cost off-policy evaluations.** This enables policy improvement through historical replay: the agent can “dream” over many alternative exploration strategies before redeploying the improved policy online.

to terminate. Executing the policy online produces a structured discovery history containing the explored branches, decision points, computational costs, and realized outcomes. As illustrated in Figure 2, this history can subsequently be treated as an *empirical replay simulator*: a grounded model of the portion of the discovery space that has already been observed.

Within this replay simulator, alternative exploration policies induce different trajectories through the recorded discovery tree. A policy may select a different subset of branches, prioritize them in a different order, issue different requests in parallel, or stop at an earlier point. Evaluating such a trajectory requires only revealing the outcomes already stored along the selected branches, rather than rerunning the underlying coding agent and evaluator. Consequently, a single expensive online discovery run can support many inexpensive evaluations of alternative exploration strategies.

3. DREAM-RSI: Recursive Self-Improvement through Evolving Worlds

As shown in Figure 1, DREAM-RSI alternates between online exploration and offline “dreaming” to improve an executable *exploration policy* that allocates discovery computation. During the online phase, the policy guides a fixed *discovery agent*, while a fixed *evaluator* scores the resulting candidates and provides diagnostic feedback. The resulting *discovery tree* serves as a *replay world* in which alternative policies can be evaluated using recorded outcomes. A fixed LLM-based *policy-development agent* uses this feedback to revise the *exploration policy* code, and the best evaluated version is deployed for the next online rollout. Only the exploration-policy code changes; the underlying models, evaluator, and execution interfaces remain fixed.

Discovery trees and the shared decision interface. A discovery tree is rooted at r , which represents the initial workspace state. Each non-root node v has exactly one *primary parent*, either the root or a previously created node. This parent identifies where the attempt in v begins: the discovery agent resumes the parent’s saved workspace and uses its accumulated observations as context to produce a new attempt. Node v preserves this inherited history and records the outcome of the new generation–evaluation attempt, including the resulting filesystem snapshot, generated artifact, evaluation diagnostics, and score s_v . Scores follow a fixed task-scoring protocol, with larger values indicating better quality.

In both online execution and offline replay, the exploration policy observes a tree $\mathcal{T}$, initially containing only the root, and selects the nodes from which to continue exploration. The eligible nodes form the set $A(\mathcal{T}) = \{r\} \cup \{v \in \mathcal{T} : v \text{ is a leaf}\}$, where leaves are determined from the currently observed tree. Let $W \geq 1$ be the number of parallel workers, each of which can execute one generation–evaluation request at a time (e.g. concurrent API calls). The exploration policy’s action is a batch $C \in A(\mathcal{T}; W)$, where $A(\mathcal{T}; W) = \{C \subseteq A(\mathcal{T}) : |C| \leq W\}$ is the feasible batch set. Each selected node specifies the starting point of one attempt, so the batch determines both where exploration continues and how many attempts are scheduled in parallel. Both the online and offline phases use this same decision interface but differ in the transition that follows a selected batch.

Online rollout. Let $t = 1, 2, \dots$ index the outer iterations, starting from an initial policy π_1 and an empty history $\mathcal{H}_0 = ()$. At iteration t , policy π_t guides a new online rollout with access to the completed discovery history $\mathcal{H}_{t-1}$. This history provides context for exploration but remains separate from the new tree being constructed. The policy code stays fixed throughout the rollout.

Let $\mathcal{T}_t^k$ denote the new discovery tree after k completed decision rounds, with $\mathcal{T}_t^0 = \{r\}$. The rollout allows at most K_1 rounds. At round $k \leq K_1$, the *exploration policy* chooses a node batch $C_t^k \in A(\mathcal{T}_t^k; W)$ and each node $v \in C_t^k$ is assigned to a worker. The discovery agent uses v ’s saved workspace and available context to produce a new candidate, and the evaluator assesses the result. These attempts run in parallel, each producing one new child of its selected parent. Attaching the completed children to the current tree yields $\mathcal{T}_t^{k+1}$, while all previously recorded nodes remain unchanged. This transition is stochastic because the discovery agent may generate different outcomes from the same starting workspace. For the next round, the newly created child becomes the selectable leaf of an extended branch, while the root remains selectable for opening further branches. The rollout ends when the policy selects an empty batch or completes K_1 decision rounds. After the rollout terminates, its final tree is recorded as $\mathcal{T}_t$ and appended to the history, giving $\mathcal{H}_t = \mathcal{H}_{t-1} \cup \{\mathcal{T}_t\}$. The method then enters the offline phase using this expanded collection of replay worlds.

Offline evaluation. During the offline phase of outer iteration t , the history $\mathcal{H}_t$ remains fixed while the method constructs and evaluates $M \geq 1$ policy versions $\pi_t^0, \dots, \pi_t^{M-1}$, starting with $\pi_t^0 = \pi_t$. Each version is evaluated separately on every historical tree $\mathcal{T}_i$, $i = 1, \dots, t$, before the next version is developed from the resulting feedback. We use m to index policy versions, i to index replay worlds, and k to count decision rounds within one policy–world evaluation. The outer index t is fixed throughout this phase and is suppressed in the notation for replay trajectories and scores.

For each policy–tree pair (m, i) , replay resets the policy’s per-rollout state and starts from $\mathcal{T}_i^{m,0} = \{r\}$. Here, $\mathcal{T}_i^{m,k} \subseteq \mathcal{T}_i$ denotes the subtree revealed after k completed rounds. The full recorded tree $\mathcal{T}_i$ remains fixed; only the portion observed by the policy evolves. At each decision, π_i^m selects a batch $C_i^{m,k} \in A(\mathcal{T}_i^{m,k}; W)$ using the revealed observations. Unlike online execution, replay returns recorded children of the selected nodes deterministically rather than generating new candidates. After the *exploration policy* takes a nonempty batch $C_i^{m,k}$, the next observed tree is $\mathcal{T}_i^{m,k+1} = \mathcal{T}_i^{m,k} \cup$

$\bigcup_{v \in C_i^{m,k}} \text{Child}(v; \mathcal{T}_i, \mathcal{T}_i^{m,k})$ where $\text{Child}(v; \mathcal{T}_i, \mathcal{T}_i^{m,k})$ denotes the node set containing unobserved children of v on tree $\mathcal{T}_i$ given the current observed tree $\mathcal{T}_i^{m,k}$. For $v \neq r$, $\text{Child}(v; \mathcal{T}_i, \mathcal{T}_i^{m,k})$ is v 's unique recorded child, if one exists. Since v is a leaf of $\mathcal{T}_i^{m,k}$, that child is still unrevealed. For $v = r$, replay returns the earliest-created child of r outside $\mathcal{T}_i^{m,k}$, opening one previously unrevealed branch. In either case, $\text{Child}(v; \mathcal{T}_i, \mathcal{T}_i^{m,k}) = \emptyset$ when no recorded continuation remains. The newly revealed nodes expose their stored observations before the policy makes its next decision.

Replay allows at most K_2 decision rounds where each nonempty batch counts as one round, and terminates when the policy selects $C_i^{m,k} = \emptyset$, the round limit $k = K_2$ is reached, or $\mathcal{T}_i^{m,k} = \mathcal{T}_i$, meaning that all recorded nodes have been revealed. Let $k_i^{m,\star} \in \{0, \dots, K_2\}$ denote the number of completed rounds at termination, yielding the final subtree $\mathcal{T}_i^{m,k_i^{m,\star}} \subseteq \mathcal{T}_i$.

Thus, replay evaluates how far to pursue each opened branch, how to group attempts into parallel batches, and when to open another branch or stop. These decisions may differ across policies, but each branch is traversed in its recorded parent-child order, and no outcomes beyond $\mathcal{T}_i$ are generated.

Replay objective. The replay objective balances discovery quality, execution cost, and parallelism. Let $N_i^m = |\mathcal{T}_i^{m,k_i^{m,\star}}| - 1$ be the number of revealed non-root nodes. Although replay itself does not execute new discovery attempts, N_i^m counts the generation-evaluation requests represented by its trajectory. For fixed coefficients $\beta_1, \beta_2 \geq 0$, the replay score is

$$V_i^m = \underbrace{\max_{v \in \mathcal{T}_i^{m,k_i^{m,\star}}} s_v}_{\text{discovery quality}} - \underbrace{\beta_1 N_i^m}_{\text{execution cost}} + \underbrace{\beta_2 \frac{N_i^m}{\max\{1, k_i^{m,\star}\}}}_{\text{parallelism bonus}}. \quad (1)$$

The first term measures the best solution quality attained during replay. The second penalizes the number of attempted generations. For a nonempty replay, the third rewards the average number of attempts executed per decision round, favoring policies that batch useful continuations rather than execute them sequentially.

Policy improvement and selection. The evaluation score of policy version π_t^m is its average replay score across the fixed history, $V^m = \frac{1}{t} \sum_{i=1}^t V_i^m$. The offline phase begins by evaluating the current policy $\pi_t^0 = \pi_t$. For each $m = 0, \dots, M-1$, the *policy-development agent* examines the replay trajectories and scores of π_t^m , together with feedback from earlier revisions, to identify successful decisions and recurring failures. It then revises the executable policy code to produce π_t^{m+1} , which is evaluated on the same t replay worlds. Replay feedback is available to the development agent between revisions.

After M revisions, the next online policy is selected from all M evaluated versions as $\pi_{t+1} = \pi_t^{m^\star}$, where $m^\star \in \arg \max_{m \in \{0, \dots, M-1\}} V^m$. Because the candidate set includes the current policy, this selection satisfies $V^{m^\star} \geq V^0$. Thus, the selected policy π_{t+1} is no worse than the current policy π_t in average replay score on the fixed history $\mathcal{H}_t$. The selected policy is then deployed online to collect $\mathcal{T}_{t+1}$, expanding the history available for the next offline improvement phase.

4. Experiments

We evaluate DREAM-RSI across three scientific discovery domains: algorithm engineering, kernel optimization and math optimization. Our primary controlled baseline is Recursive Fixed Exploration,

which uses the same underlying discovery setting and initialization but keeps the exploration policy fixed across recursive discovery rounds. We additionally compare against task-specific domain baselines.

Across all tasks, DREAM-RSI and Recursive Fixed Exploration use the same discovery agent, evaluator, initialization, and resource constraints. Both methods start from the same manually designed exploration policy. This exploration policy follows a simple *parallel refining* strategy: it launches multiple independent exploration workspaces in parallel, with each workspace maintaining its own local discovery trajectory and repeatedly refining its current candidate based on the history accumulated within that workspace. The two methods therefore follow the same exploration policy in the first discovery round. In subsequent rounds, while Recursive Fixed Exploration keeps its exploration policy static, DREAM-RSI progressively refines the policy by dreaming over a replay simulator conditioned on accumulated global discovery history, subsequently deploying the updated policy in each new round. The discovery cost is quantified by the total cumulative number of discovery-agent calls.

Specifically, we evaluate Gemini-3.1 Pro and Gemini-3.7-Flash across multiple recursive discovery rounds via the Gemini CLI ². Under Recursive Fixed Exploration, each round for Gemini-3.1 Pro executes 10 parallel workspaces with up to 11 refinement steps ($10 \times 11 = 110$ discovery-agent calls), whereas Gemini-3.7-Flash operates 32 parallel workspaces with up to 20 refinement steps ($32 \times 20 = 640$ calls). DREAM-RSI maintains identical per-round budgets, aligning with the baseline in Round 1 while progressively updating its policy in subsequent rounds. Further details on recursive rounds, task setups, resource budgets, and evaluation protocols follow below.

4.1. Algorithm Engineering

In this task, we consider **Lasso Regularization Path** as our algorithm-engineering task, a fundamental computational primitive in high-dimensional statistics that is widely used in model selection and cross-validation across domains such as genomics and finance. We follow the benchmark setting of SimpleTES (Ye et al., 2026), where the goal is to discover efficient implementations of the complete Lasso regularization path while preserving numerical correctness. During discovery, we use the same 17 synthetic instances as SimpleTES, which cover diverse problem regimes in terms of dimensionality, sparsity, feature correlation, and active-set structure. To evaluate whether the discovered algorithms generalize beyond the search distribution, we additionally evaluate them on six held-out downstream datasets spanning both biological and non-biological domains.


Baselines and Setup. We compare against standard Lasso solvers **sklearn** (Pedregosa et al., 2011) and **glmnet** (Friedman et al., 2010), as well as **SimpleTES** (Ye et al., 2026), which uses GPT-OSS-120B with a reported budget of 51,200 generations. We additionally include Recursive Fixed Exploration as our controlled baseline. Specifically, we run both Recursive Fixed Exploration and DREAM-RSI for 5 rounds.

Main Results. Figure 3(a) summarizes the Lasso discovery results. Across both discovery-agent backbones, DREAM-RSI achieves a better downstream quality–compute trade-off than Recursive Fixed Exploration. With Gemini-3.1 Pro, it reduces the average runtime across the six held-out datasets from 3587.1 ms to 2931.0 ms while using only 317 discovery-agent calls, compared with 550 calls for fixed exploration. With Gemini-3.7-Flash, DREAM-RSI further reduces the average runtime from 2516.7 ms to 2350.6 ms using 1879 calls instead of 3200. Despite using substantially less discovery compute,

²<https://geminicli.com/>

Method	Model	Compute	Non-biological		Biological				Avg.
			Gisette	RCV1	DNA	Leukemia	Colon	Duke Breast	
Previous solvers									
sklearn	–	–	11275.2	252881.7	93.8	227.2	229.8	374.0	44180.3
glmnet	–	–	9063.6	73072.8	351.9	45.0	24.2	47.7	13767.5
SimpleTES	gpt-oss-120b	51,200	3141.9	19625.6	15.9	15.5	11.6	18.1	3804.8
SimpleTES †	gpt-oss-120b	51,200	8651.0	41143.1	37.6	28.2	19.5	31.1	8318.4
Our System									
Recursive Fixed Exploration	Gemini-3.1-Pro	550	1861.8	19550.1	41.5	26.1	14.5	28.4	3587.1
	Gemini-3.7-Flash	3200	1133.1	13873.0	29.8	24.1	15.7	24.4	2516.7
DREAM-RSI	Gemini-3.1-Pro	317	2841.0	14616.0	49.9	30.2	16.4	32.5	2931.0
	Gemini-3.7-Flash	1879	1091.9	12923.4	31.4	21.0	12.2	23.6	2350.6

(a) Final performance.



(b) Recursive Discovery Dynamics.

Figure 3 | **Lasso regularization-path discovery results.** (a) Final wall-clock runtime on six held-out downstream tasks; lower is better. **Compute** denotes the cumulative number of discovery-agent calls. (b) **Recursive discovery dynamics.** Average downstream runtime across six held-out tasks versus cumulative discovery compute for Gemini-3.1-Pro and Gemini-3.7-Flash. Numbers next to markers denote recursive rounds (iterations). Lower is better.

the resulting solvers also outperform the standard sklearn and glmnet implementations on all six held-out datasets. Compared with SimpleTES, which uses 51,200 generations, DREAM-RSI achieves lower average downstream runtime with roughly two orders of magnitude fewer discovery-agent calls. Notably, the program discovered by Gemini-3.1-Pro appears particularly well suited to large-scale matrices such as RCV1. In contrast, Gemini-3.7-Flash discovers a more general-purpose program that performs consistently across different problem scales.

Recursive Discovery Dynamics. Figure 3(b) illustrates the trajectory of downstream performance across recursive discovery rounds relative to cumulative discovery compute. By design, both methods share identical search behavior in the initial round. In subsequent rounds, Recursive Fixed Exploration maintains a static exploration policy, whereas DREAM-RSI progressively refines and redeploys its policy via dreaming over accumulated discovery history. Consequently, the two trajectories diverge markedly: DREAM-RSI consistently achieves superior downstream performance while requiring substantially lower cumulative compute across both Gemini-3.1-Pro and Gemini-3.7-Flash.

Table 1 | Performance comparison on mathematical discovery tasks. Higher is better for **Sum Diff** and **Circle Packing**, while lower is better for **Auto Correlation**. Best results are shown in **bold**.

Method	LLM	Sum Diff ($\uparrow$)	Auto Correlation ($\downarrow$)	Circle Packing ($\uparrow$)
AlphaEvolve	Gemini-2.0 Pro + Flash	–	1.455700	2.635862
AlphaEvolveV2	Gemini-2.0 Pro + Flash	1.121936	–	2.635983
OpenEvolve	-	–	1.460000	-
CodeEvolve	-	–	–	2.635980
ShinkaEvolve	Mixed	–	1.457800	2.635982
TTS-Discovery	Qwen3-8B	–	–	2.635983
ThetaEvolve	Distilled-Qwen3-8B	–	1.493000	2.635983
EvoX	Gemini-3.0-Pro	–	1.458900	2.635900
SimpleTES	GPT-OSS-120B	1.143975	1.453675	2.635983
<i>Our System</i>				
Recursive Fixed Exploration	Gemini-3.1-Pro	1.144047	1.456001	2.635983
DREAM-RSI	Gemini-3.1-Pro	1.145427	1.456375	2.635983

Discovered Solver Analysis. We further analyze the discovered solver, with its implementation provided in the Appendix C. Unlike SimpleTES, which switches between LARS and coordinate descent according to problem dimensions, the discovered solver introduces adaptivity within the active-set optimization itself. It combines strong-rule screening with Cauchy–Schwarz-based KKT pruning, selectively recomputing exact gradients only when the bound cannot certify a feature and falling back to a full refresh when pruning becomes ineffective. This adaptive verification scheme is further integrated with efficient active-set bookkeeping, lazy Gram-matrix construction, and hardware-aware implementation.

4.2. Mathematics Optimization

We further evaluate DREAM-RSI on three mathematical discovery tasks spanning discrete combinatorial optimization, geometric optimization, and functional optimization: the **Sum-Difference Problem**, **Circle Packing**, and **Autocorrelation Inequalities**. The goal of these problems is to discover high-quality solutions that optimize task-specific mathematical objectives under their respective constraints. Formal definitions of the three tasks are provided in Appendix.

We use Gemini-3.1 Pro via the Gemini CLI as the discovery agent for both Recursive Fixed Exploration and DREAM-RSI for 10 rounds. For each task, the agent iteratively proposes and evaluates candidate constructions or optimization procedures according to the task-specific objective. We compare against a broad set of existing automated discovery systems, including AlphaEvolve (Novikov et al., 2025), AlphaEvolveV2 (Georgiev et al., 2025), OpenEvolve (Sharma, 2025), CodeEvolve (Assumpção et al., 2025), ShinkaEvolve (Lange et al., 2026), TTS-Discovery (Yuksekgonul et al., 2026), ThetaEvolve (Wang et al., 2025), EvoX (Liu et al., 2026a), and SimpleTES (Ye et al., 2026).

Results. Table 1 summarizes the results across the three mathematical discovery tasks. DREAM-RSI achieves a Sum-Difference score of 1.145427, outperforming SimpleTES and Recursive Fixed Exploration. On Circle Packing, it reaches 2.635983, matching the strongest reported result among the compared methods. For Autocorrelation, DREAM-RSI obtains 1.456375, remaining competitive with existing discovery systems. Notably, SimpleTES achieves state-of-the-art performance on Autocorre-




Figure 4 | **GPU kernel engineering results.** Discovery performance of DREAM-RSI and Recursive Fixed Exploration as a function of the number of generations. On VGG16 and LayerNorm, DREAM-RSI reaches comparable performance with 2.43× and 1.79× fewer generations, respectively. On ConvDiv and ConvMax, it achieves 2.09× and 1.44× higher performance under comparable discovery budgets. Higher is better for all tasks.

lation Inequalities, but requires 51,200 generations, significantly more than the fewer than 1,000 generations used by our approach. Overall, these results show that our DREAM-RSI generalize well on mathematics optimization.

4.3. Kernel Engineering

We further evaluate DREAM-RSI on **GPU kernel engineering**, where the goal is to automatically discover high-performance implementations of kernels while preserving numerical correctness. Unlike mathematical discovery, kernel engineering requires reasoning jointly about algorithmic structure, memory access, parallelization, and hardware-specific optimizations, providing a substantially different testbed for evaluating whether our DREAM-RSI generalizes across discovery domains.

We consider four representative kernel-engineering tasks from KernelBench (Ouyang et al., 2025): **VGG16**, **LayerNorm**, **ConvDiv**, and **ConvMax**. Candidate implementations are evaluated by their execution performance, measured as inverse runtime (1/ms), subject to correctness checks against the reference implementation. We use Gemini-3.1 Pro as the coding agent and compare DREAM-RSI with Recursive Fixed Exploration under the same evaluation protocol and initialization.

Results. Figure 4 shows the discovery trajectories as the number of generations increases. On VGG16 and LayerNorm, DREAM-RSI reaches comparable final performance using 2.43× and 1.79× fewer generations, respectively. On ConvDiv and ConvMax, under comparable discovery budgets,

DREAM-RSI achieves $2.09\times$ and $1.44\times$ higher performance, respectively. These results show that adapting the exploration policy across recursive rounds can improve the efficiency and effectiveness of long-horizon discovery.

5. Further Analysis

5.1. Analysis of Historical Inductive Biases in Long-Horizon Discovery

We further investigate how the nature of the historical inductive bias affects long-horizon discovery. A natural alternative for utilizing history is to abstract prior trajectories into high-level directional insights, which are directly injected into the prompt as explicit semantic guidance for subsequent rounds. To evaluate the efficacy of this prompt-level semantic guidance, we apply it to both Recursive Fixed Exploration and DREAM-RSI. As illustrated in Figure 5, explicit directional guidance consistently underperforms its unguided counterpart across both paradigms under equivalent discovery budgets. These results suggest that in long-horizon discovery—where multiple parallel threads are deployed for exploration—imposing strong semantic inductive biases regarding future search directions tends to over-constrain the search space and impede diverse exploration.




Figure 5 | Discovery performance on ConvDiv. Using history as an interactive replay simulator outperforms using it only as guidance.

5.2. Analysis of Evolution of Exploration Behavior

Figure 6 illustrates how the learned exploration policy evolves across recursive rounds on ConvDiv. As shown, the exploration policy exhibits a clear adaptive pattern: as performance improves, it initially conserves discovery compute (e.g., reducing the number of evaluated attempts from 110 to 50). When progress subsequently plateaus, it increases exploration effort again, coinciding with further performance gains.




Figure 6 | Evolution of exploration behavior on ConvDiv. (a) Round-best performance across recursive execution rounds. (b) The number of evaluated attempts in each round.

6. Related Work

AI-Driven Scientific and Algorithmic Discovery. LLM-based discovery systems iteratively generate, evaluate, and refine candidate solutions using prior artifacts and feedback, as in AlphaEvolve (Novikov et al., 2025), OpenEvolve (Sharma, 2025), CodeEvolve (Assumpção et al., 2025), ShinkaEvolve (Lange et al., 2026), PACEvolve (Yan et al., 2026b), DeltaEvolve (Jiang et al., 2026) and MLEvolve Du et al.

(2026). More recent work emphasizes the importance of exploration itself: SkyDiscover provides adaptive discovery infrastructure (Liu et al., 2026b), SwarmResearch dynamically orchestrates multiple search branches (Virk et al., 2026), and EvoX (Liu et al., 2026a) explicitly optimizes search strategies rather than only candidate solutions. This shift makes exploration a meta-level optimization problem, but useful supervision for exploration strategies is expensive and delayed because their quality often becomes apparent only after long discovery rollouts.

Self-Evolving Agents. A broader line of work studies agents that improve their own components during interaction. Prior methods evolve model weights (Huang et al., 2026a,c), agent harnesses (Lee et al., 2026; Zhang et al., 2026b), contexts (Zhang et al., 2026d), skills (Ouyang et al., 2026a; Wu et al., 2026b; Zhang et al., 2026a), model behavior through test-time learning (Wang et al., 2025; Wu et al., 2026a; Yan et al., 2026a; Yuksekogonul et al., 2026), rubrics (Xiong et al., 2026), environments (Huang et al., 2026b) and other applications (Dai et al., 2026). Most operate at the *object level*, improving components used for task execution or reasoning. Recent work has begun to optimize meta-level mechanisms, including search strategies and self-improvement procedures (Kim et al., 2026; Liu et al., 2026a; Wang et al., 2026; Yan et al., 2026a; Zhang et al., 2026c). However, such meta-level strategies are difficult to improve because their quality is often revealed only after costly long-horizon rollouts. DREAM-RSI makes this meta-level optimization recursive and off-policy by turning accumulated discovery history into replay simulators, allowing exploration controllers to be repeatedly evaluated, improved, and redeployed without rerunning the underlying discovery process.

Memory, History, and Experience Reuse. Prior work reuses agent experience as search history, context, memory, reusable skills, or training signals. DeltaEvolve structures evolutionary history through semantic deltas (Jiang et al., 2026); SwarmResearch and MLEvolve use cross-branch or retrospective information to guide subsequent search (Du et al., 2026; Virk et al., 2026); and other work improves how agents access and retain experience through evolving contexts, broader harness state, libraries, or skills (Lee et al., 2026; Ouyang et al., 2026a; Xu et al., 2026; Zhang et al., 2026d). We take a different view: rather than using exploration history only as context or memory for the next decision, we organize it as a *replay simulator* in which many alternative exploration controllers can be evaluated cheaply. This turns previously collected discovery experience into reusable feedback for meta-level optimization, alleviating the scarcity and high cost of training signals for improving exploration strategies.

7. Conclusion

We presented DREAM-RSI, a framework for recursive self-improvement of exploration in recursive self improvement. By converting accumulated discovery history from static context into an active, replayable simulator, DREAM-RSI addresses the core bottleneck of meta-optimization: delayed and expensive feedback, which is especially severe in long-horizon discovery settings. By ‘dreaming’ within replay simulators constructed from historical discovery trees, Dream-RSI evaluates candidate exploration policies rapidly and at negligible execution cost. The improved policies are then redeployed online to drive further discovery and expand the simulator pool, closing the recursive self-improvement loop. Across algorithm engineering, mathematical optimization, and GPU kernel engineering, DREAM-RSI achieves competitive or improved discovery quality while substantially reducing discovery cost in several settings.

References

- Anthropic. Learning more about claude’s mathematical capabilities. <https://www.anthropic.com/research/riemann-zeta>, Aug. 2026. Accessed: 2026-08-13.
- H. Assumpção, D. Ferreira, L. Campos, and F. Murai. Codeevolve: an open source evolutionary coding agent for algorithmic discovery and optimization. *arXiv preprint arXiv:2510.14150*, 2025.
- S. Cao, Z. Mao, J. E. Gonzalez, and I. Stoica. K-search: Llm kernel generation via co-evolving intrinsic world model. *arXiv preprint arXiv:2602.19128*, 2026.
- R. Dai, K. Huang, C. Kang, and C. Liao. It takes two to match: Co-evolving generative retriever with reinforcement learning. *arXiv preprint arXiv:2609.00638*, 2026.
- S. Du, X. Yan, J. Shi, Z. Cao, S. Feng, Z. Liang, B. Sun, T. Peng, Y. Zhou, X. Li, et al. Mlevolve: A self-evolving framework for automated machine learning algorithm discovery. *arXiv preprint arXiv:2606.06473*, 2026.
- J. H. Friedman, T. Hastie, and R. Tibshirani. Regularization paths for generalized linear models via coordinate descent. *Journal of statistical software*, 33:1–22, 2010.
- B. Georgiev, J. Gómez-Serrano, T. Tao, and A. Z. Wagner. Mathematical exploration and discovery at scale. *arXiv preprint arXiv:2511.02864*, 2025.
- S. Gupta, J. Davidson, S. Levine, R. Sukthankar, and J. Malik. Cognitive mapping and planning for visual navigation. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 2616–2625, 2017.
- D. Ha and J. Schmidhuber. World models. *arXiv preprint arXiv:1803.10122*, 2(3):440, 2018.
- D. Hafner, T. Lillicrap, J. Ba, and M. Norouzi. Dream to control: Learning behaviors by latent imagination. *arXiv preprint arXiv:1912.01603*, 2019.
- D. Hafner, T. Lillicrap, M. Norouzi, and J. Ba. Mastering atari with discrete world models. *arXiv preprint arXiv:2010.02193*, 2020.
- D. Hafner, J. Pasukonis, J. Ba, and T. Lillicrap. Mastering diverse domains through world models. *arXiv preprint arXiv:2301.04104*, 2023.
- D. Hafner, W. Yan, and T. Lillicrap. Training agents inside of scalable world models. *arXiv preprint arXiv:2509.24527*, 2025.
- Y. Hu, S. Liu, Y. Yue, G. Zhang, B. Liu, F. Zhu, J. Lin, H. Guo, S. Dou, Z. Xi, et al. Memory in the age of ai agents. *arXiv preprint arXiv:2512.13564*, 2025.
- C. Huang, H. Liu, T. Zheng, R. Dai, L. Huang, J. Li, Z. Li, Z. Wei, Y. Meng, and J. Huang. G-zero: Self-play for open-ended generation from zero data. *arXiv preprint arXiv:2605.09959*, 2026a.
- C. Huang, Z. Wang, R. Han, J. Yan, Y. Chen, Z. CuiZhu, K. Jiang, P. Xia, H. Yu, Y. Zhuang, et al. Envharness: Awakening static worlds for agent learning. *arXiv preprint arXiv:2608.19880*, 2026b.
- C. Huang, W. Yu, X. Wang, H. Zhang, Z. Li, R. Li, J. Huang, H. Mi, and D. Yu. R-zero: Self-evolving reasoning llm from zero data. In *International Conference on Learning Representations*, volume 2026, pages 130770–130790, 2026c.

- J. Jaber and O. Jaber. Autokernel: Autonomous gpu kernel optimization via iterative agent-driven search. *arXiv preprint arXiv:2603.21331*, 2026.
- J. Jiang, T. Ding, and Z. Zhu. Deltaevolve: Accelerating scientific discovery through momentum-driven evolution. *arXiv preprint arXiv:2602.02919*, 2026.
- Z. M. Kim, Y.-J. Lee, S. Jwa, and D. Kang. Metaⁿ: Recursive self-improvement through emergent depth. *arXiv preprint arXiv:2608.24735*, 2026.
- R. Lange, Y. Imajuku, and E. Cetin. Shinkaevolve: Towards open-ended and sample-efficient program evolution. In *International Conference on Learning Representations*, volume 2026, pages 74026–74078, 2026.
- Y. Lee, R. Nair, Q. Zhang, K. Lee, O. Khattab, and C. Finn. Meta-harness: End-to-end optimization of model harnesses. *arXiv preprint arXiv:2603.28052*, 2026.
- S. Liu, S. Agarwal, M. Maheswaran, M. Cemri, Z. Li, Q. Mang, A. Naren, E. Boneh, A. Cheng, M. Z. Pan, et al. Evox: Meta-evolution for automated discovery. *arXiv preprint arXiv:2602.23413*, 2026a.
- S. Liu, M. Cemri, S. Agarwal, A. Krentsel, A. Naren, Q. Mang, Z. Li, A. Gupta, M. Maheswaran, A. Cheng, M. Pan, E. Boneh, K. Ramchandran, K. Sen, M. Zaharia, A. G. Dimakis, and I. Stoica. Skydiscover: A flexible, adaptive framework for ai-driven scientific and algorithmic discovery. In *Proceedings of the ACM Conference on AI and Agent Systems*, CAIS ’26, pages 1223–1227. Association for Computing Machinery, 2026b. doi: 10.1145/3786335.3813221. URL <https://doi.org/10.1145/3786335.3813221>.
- S. Liu, Z. Lin, Y. Zhang, Y. Ren, Y. Wu, Y. Li, Z. Wang, Z. Fu, and J. Ye. The path to recursive self-improving agents: Foundation, framework, and future directions. *Preprints*, August 2026c. doi: 10.20944/preprints202608.0051.v1. URL <https://doi.org/10.20944/preprints202608.0051.v1>.
- T. M. Moerland, J. Broekens, A. Plaat, and C. M. Jonker. Model-based reinforcement learning: A survey. *Foundations and Trends in Machine Learning*, 16(1):1–118, 2023.
- A. Novikov, N. Vű, M. Eisenberger, E. Dupont, P.-S. Huang, A. Z. Wagner, S. Shirobokov, B. Kozlovskii, F. J. Ruiz, A. Mehrabian, et al. Alphaevolve: A coding agent for scientific and algorithmic discovery. *arXiv preprint arXiv:2506.13131*, 2025.
- OpenAI. On the navier–stokes millennium prize problem. <https://openai.com/index/navier-stokes-solution/>, Sept. 2026. Accessed: 2026-09-10.
- A. Ouyang, S. Guo, S. Arora, A. L. Zhang, W. Hu, C. Ré, and A. Mirhoseini. Kernelbench: Can llms write efficient gpu kernels? *arXiv preprint arXiv:2502.10517*, 2025.
- S. Ouyang, J. Yan, Y. Chen, R. Han, Z. Wang, B. D. Mishra, R. Meng, C.-L. Li, Y. Jiao, K. Zha, et al. Skillos: Learning skill curation for self-evolving agents. *arXiv preprint arXiv:2605.06614*, 2026a.
- S. Ouyang, J. Yan, I. Hsu, Y. Chen, K. Jiang, Z. Wang, R. Han, L. Le, S. Daruki, X. Tang, et al. Reasoningbank: Scaling agent self-evolving with reasoning memory. In *International Conference on Learning Representations*, volume 2026, pages 94327–94354, 2026b.
- F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, et al. Scikit-learn: Machine learning in python. *the Journal of machine Learning research*, 12:2825–2830, 2011.

- B. Romera-Paredes, M. Barekatain, A. Novikov, M. Balog, M. P. Kumar, E. Dupont, F. J. Ruiz, J. S. Ellenberg, P. Wang, O. Fawzi, et al. Mathematical discoveries from program search with large language models. *Nature*, 625(7995):468–475, 2024.
- A. Sharma. Openevolve: an open-source evolutionary coding agent, 2025. URL <https://github.com/algorithmicsuperintelligence/openevolve>.
- R. S. Sutton. Integrated architectures for learning, planning, and reacting based on approximating dynamic programming. In B. Porter and R. Mooney, editors, *Machine Learning Proceedings 1990*, pages 216–224. Morgan Kaufmann, San Francisco (CA), 1990. ISBN 978-1-55860-141-3. doi: <https://doi.org/10.1016/B978-1-55860-141-3.50030-4>. URL <https://www.sciencedirect.com/science/article/pii/B9781558601413500304>.
- Y. Virk, Z. Edds, C. S. Xia, and L. Zhang. Swarmresearch: Orchestrating coding agents for open-ended discovery. *arXiv preprint arXiv:2607.02807*, 2026.
- Y. Wang, S.-R. Su, Z. Zeng, E. Xu, L. Ren, X. Yang, Z. Huang, X. He, L. Ma, B. Peng, et al. Thetaevolve: Test-time learning on open problems. *arXiv preprint arXiv:2511.23473*, 2025.
- Z. Wang, M. Yan, J. Bi, S. Yan, V. Tresp, and Y. Ma. Metaskill-evolve: Recursive self-improvement of llm agents via two-timescale meta-skill evolution. *arXiv preprint arXiv:2607.05297*, 2026.
- S. Wu, C. Qian, X. Chen, and H. Ji. Teaching llms to self-evolve: Cultivating core meta-skills with reinforcement learning. *arXiv preprint arXiv:2607.21971*, 2026a.
- X. Wu, Y. Zhuan, R. Wei, H. Chen, D. Bai, J. Liu, X. Wang, X. Wang, L. Wang, and X. Cheng. Agenticrectune: Multi-agent with self-evolving skillhub for recommendation system optimization. *arXiv preprint arXiv:2604.26969*, 2026b.
- T. Xiong, Z. Yang, X. Wang, C.-C. Lin, R. Ma, K. Lin, Z. Wang, L. Li, C. Liu, R. Chen, et al. Rubrics as visual-repair context for self-evolving ui-to-code generation. *arXiv preprint arXiv:2608.24138*, 2026.
- W. Xu, A. Sordoni, C. Singh, Z. Gero, M. Galley, X. Yuan, and J. Gao. Test-time learning with an evolving library. *arXiv preprint arXiv:2605.14477*, 2026.
- M. Yan, B. Peng, B. Coleman, Z. Chen, Z. Xie, S. Chen, Z. He, N. Sachdeva, W. Wang, E. H. Chi, et al. Paceevolve++: Improving test-time learning for evolutionary search agents. *arXiv preprint arXiv:2605.07039*, 2026a.
- M. Yan, B. Peng, B. Coleman, Z. Chen, Z. Xie, S. Chen, Z. He, N. Sachdeva, I. Ye, W. Wang, et al. Paceevolve: Enabling long-horizon progress-aware consistent evolution. *arXiv preprint arXiv:2601.10657*, 2026b.
- H. Ye, H. Lin, J. Tang, Y. Luo, C. Yang, C. Su, R. Thapa, R. Yang, R. Liu, Z. Li, et al. Evaluation-driven scaling for scientific discovery. *arXiv preprint arXiv:2604.19341*, 2026.
- M. Yuksekgonul, D. Kocejka, X. Li, F. Bianchi, J. McCaleb, X. Wang, J. Kautz, Y. Choi, J. Zou, C. Guestrin, et al. Learning to discover at test time. *arXiv preprint arXiv:2601.16175*, 2026.
- H. Zhang, S. Fan, H. P. Zou, Y. Chen, Z. Wang, J. Zhou, C. Li, W.-C. Huang, Y. Yao, K. Zheng, et al. Co-evoskills: Self-evolving agent skills via co-evolutionary verification. *arXiv preprint arXiv:2604.01687*, 2026a.

- J. Zhang, S. Hu, C. Lu, R. Lange, and J. Clune. Darwin gödel machine: open-ended evolution of self-improving agents. In *International Conference on Learning Representations*, volume 2026, pages 104223–104294, 2026b.
- J. Zhang, B. Zhao, W. Yang, J. Foerster, J. Clune, M. Jiang, S. Devlin, and T. Shavrina. Hyperagents. *arXiv preprint arXiv:2603.19461*, 2026c.
- Q. Zhang, C. Hu, S. Upasani, B. Ma, F. Hong, V. Kamanuru, J. Rainton, C. Wu, M. Ji, H. Li, et al. Agentic context engineering: Evolving contexts for self-improving language models. In *International Conference on Learning Representations*, volume 2026, pages 86069–86100, 2026d.
- T. Zheng, H. Liu, C. Huang, H. Bao, S. Zhang, R. Liu, R. Dai, R. Chen, C. Liu, T. Xiong, et al. Llms improving llms: Agentic discovery for test-time scaling. *arXiv preprint arXiv:2605.08083*, 2026a.
- T. Zheng, H. Zhang, W. Yu, X. Wang, H. Xing, R. Dai, R. Liu, H. Bao, C. Huang, H. Huang, et al. Parallel-r1: Towards parallel thinking via reinforcement learning. In *International Conference on Learning Representations*, volume 2026, pages 121144–121166, 2026b.

A. Detailed Task Description

Problem 1 (Lasso Regularization Path)

Given a feature matrix $X \in \mathbb{R}^{n \times p}$, response $y \in \mathbb{R}^n$, and a decreasing sequence $\lambda_1 > \dots > \lambda_K$, define

$$F_k(w) := \frac{1}{2n} \|y - Xw\|_2^2 + \lambda_k \|w\|_1, \quad w_k^* \in \arg \min_{w \in \mathbb{R}^p} F_k(w).$$

A candidate solver returns approximate coefficients $\tilde{W} = (\tilde{w}_1, \dots, \tilde{w}_K)$. It passes the benchmark's objective-value check if

$$F_k(\tilde{w}_k) \leq F_k(w_{k,\text{slearn}}) + 10^{-6}, \quad \text{for every } k,$$

where correctness is checked on fresh instances distinct from the timing instances. If any required check fails, the search score is zero.

Otherwise, letting $\mathcal{I}$ denote the timing instances and t_i the time required to compute the complete regularization path on instance i , the search score is

$$R_{\text{search}} = \left(\prod_{i \in \mathcal{I}} t_i \right)^{-1/|\mathcal{I}|}.$$

Problem 2 (Sum-Difference Problem)

The Sum-Difference Problem asks for a finite set $A \subset \mathbb{Z}$ whose normalized sumset is large relative to its normalized difference set. The objective is

$$\Gamma(A) := \frac{\log(|A + A|/|A|)}{\log(|A - A|/|A|)},$$

where

$$A + A := \{a + a' : a, a' \in A\}, \quad A - A := \{a - a' : a, a' \in A\}.$$

The goal is to construct a finite set $A \subset \mathbb{Z}$ that maximizes $\Gamma(A)$.

Problem 3 (Circle Packing in a Unit Square)

For $n \in \{26, 32\}$, the circle-packing task asks for centers $(x_i, y_i) \in [0, 1]^2$ and radii $r_i \geq 0$ such that every circle lies inside the unit square and no two circles overlap:

$$\begin{aligned} r_i &\leq x_i \leq 1 - r_i, & r_i &\leq y_i \leq 1 - r_i, \\ (x_i - x_j)^2 + (y_i - y_j)^2 &\geq (r_i + r_j)^2, & 1 &\leq i < j \leq n. \end{aligned}$$

The objective is to maximize

$$\sum_i r_i.$$

Problem 4 (Autocorrelation Inequalities)

For an integrable function $f : \mathbb{R} \rightarrow \mathbb{R}$, define the autoconvolution

$$(f * f)(t) := \int_{\mathbb{R}} f(t-x)f(x) dx, \quad t \in \left[-\frac{1}{2}, \frac{1}{2}\right].$$

First Autocorrelation Inequality. Find a non-negative integrable function $f : \mathbb{R} \rightarrow \mathbb{R}$ supported on $[-\frac{1}{4}, \frac{1}{4}]$ such that

$$\int_{-1/4}^{1/4} f(x) dx = 1.$$

The objective is to minimize

$$\Phi_1(f) := \max_{t \in [-1/2, 1/2]} (f * f)(t).$$

Second Autocorrelation Inequality. Find a non-negative integrable function $f : \mathbb{R} \rightarrow \mathbb{R}$ supported on $[-\frac{1}{4}, \frac{1}{4}]$ such that

$$\int_{-1/4}^{1/4} f(x) dx = 1.$$

The objective is to maximize

$$\Phi_2(f) := \frac{\|f * f\|_2^2}{\|f * f\|_1 \|f * f\|_\infty}.$$

Third Autocorrelation Inequality. Find an integrable (possibly signed) function $f : \mathbb{R} \rightarrow \mathbb{R}$ supported on $[-\frac{1}{4}, \frac{1}{4}]$ such that

$$\int_{-1/4}^{1/4} f(x) dx = 1.$$

The objective is to minimize

$$\Phi_3(f) := \max_{t \in [-1/2, 1/2]} |(f * f)(t)|.$$

B. Prompts

For reproducibility, we provide the prompts used for online exploration and replay-based exploration-policy improvement. Variables enclosed by dollar signs or braces are instantiated by the runtime system before execution.

B.1. Exploration Prompt

The following prompt is used to guide the discovery agent during online exploration. It requires the agent to inspect the complete available discovery history before proposing a new solution, explicitly reason about both successful and failed attempts, and avoid repeatedly exploiting a locally saturated direction.

Listing 1 | Prompt used for online exploration.

```

1 You must read every historical proposal before proposing or implementing a new solution.
2
3 $direction_guidance
4
5 Variables ('$node_dir', '$history_dir', '$baseline_dir', '$eval_program', '$problem_file') are
   filled in by the calling system. '$node_dir' is your own attempt directory -- exclude it when
   scanning sibling 'attempt_*/' dirs.
6
7 ## 1. Read the complete history first
8
```

```

9 Before proposing anything, read every 'proposal.md' under sibling 'attempt_*/' dirs, '$history_dir',
  and '$baseline_dir' in full -- not a sample, not just recent cycles or the current branch.
  For each, read its matching 'eval/score.json' (and 'error.txt' if it failed). Trust the
  measured result over what the proposal claims about itself.
10
11 ## 2. Learn from both successes and failures
12
13 For every past attempt, note the mechanism and how it did. For failures, figure out *why*: a
  flawed core idea, or a good idea let down by a bug, bad parameters, or an implementation slip
  ? Don't repeat the former. The latter is worth retrying -- but only once you've actually
  located the bug in the code (not just guessed from the proposal), and only with a specific
  fix in hand.
14
15 ## 3. Don't converge into a local optimum
16
17 Look at the shape of what's been tried. If most attempts cluster around small variations of one
  mechanism with flattening returns, that's a local optimum - resist proposing another small
  tweak there. Deliberately favor a structurally different mechanism or an untried combination
  over a safer marginal refinement. Exploration diversity matters as much as the next
  incremental gain.
18
19 ## 4. Propose and implement
20
21 The new idea must be a genuinely new mechanism, a new combination of previously-successful pieces,
  or a targeted fix to a specific bug found in step 2 - never a repeat or rename of something
  already tried. Implement it in '$eval_program'. Don't claim it compiles, is correct, or beats
  SOTA until it's actually evaluated.
22
23 ## Files
24
25 Write only '$node_dir/proposal.md' (mechanism, evidence from history, why it's not a repeat,
  expected benefit/risk) and '$node_dir/$eval_program'. Everything else is read-only.
26
27 ## Note:
28     Never execute pkill, kill, killall, or terminate unrelated processes.

```

B.2. Replay-Based Policy Improvement Prompt

The following prompt is used by the controller-development agent during historical replay. The agent modifies the exploration policy using feedback obtained from replay over previously collected discovery trajectories while remaining restricted to prefix-observable information.

Listing 2 | Prompt used for replay-based improvement of the exploration policy.

```

1 You are improving one **prefix-only exploration policy**. Edit only
2 '{method_file}' and implement 'OptimalPolicy.solve(self, question, budget=None)'.
3 Do not solve the scientific task and do not edit any other program.
4
5 ## Objective: quality, work, and parallelism
6
7 The environment is a frozen, irregular branch×attempt grid. A policy opens a root
8 or refines the next cell of an already-open branch. Each revealed cell costs one
9 probe. The policy sees only the cells it has revealed so far; unrevealed scores are
10 unknown.
11
12 The evaluator sweeps your single 'beta' knob and ranks the resulting curve by:
13
14     pareto.reward = pareto.auc - lambda * parallel_penalty
15
16 'pareto.auc' rewards reaching high per-trace attainment with few **total probes**.
17 'parallel_penalty' is the mean of
18 'effective_sequential_rounds / total_probes' over the sweep. For a batch of size
19 'k' with 'W = question.max_parallelism' workers, it costs one decision round and
20 'ceil(k / W)' effective sequential rounds. A serial policy has penalty near 1;
21 useful full batches approach '1/W'. Therefore choose only promising probes, but
22 batch independent promising probes whenever possible.
23

```

```

24 A local implementation failure does not by itself prove that its parent direction
25 is poor. Weigh recovery value against new roots and ordinary refinements while
26 keeping batches parallel.
27
28 ## API
29
30     question.reset()
31     question.observed() -> dict[str, Observation]    # revealed prefix only
32     question.legal_actions() -> list[str]           # roots + opened-branch frontiers
33     question.legal_roots() -> list[str]             # unopened roots only
34     question.opened_branches() -> list[int]
35     question.meta(cell_id) -> CellMeta              # .branch .attempt .parent_id .seq .tags
36     question.probe_batch(cells, on_reveal=...) -> list[Observation]
37     question.baseline_score
38     question.max_parallelism
39
40     'Observation' supplies 'branch', 'attempt', 'score', 'evaluated', 'valid',
41     'fail_class', 'error', 'delta_vs_baseline', 'delta_vs_parent', 'n_valid', and
42     'n_total'.
43 Use the helpers in 'see.policy.observation_signal' when useful:
44 'branch_promising', 'branch_failed_hard', 'probe_improved_vs_parent', and
45 'probe_improved_vs_baseline'.
46
47 **Success semantics:** an evaluated observation with 'error is None' and
48 'fail_class == "ok"' is a successful evaluation, even when 'valid == False' or
49 'n_valid'/'n_total' are unavailable. Never label it repairable solely because
50 'valid' is false. A *successful anchor* below means the best historical score
51 from such a successful evaluation.
52
53 Do **not** use 'question.best_so_far' or 'question.budget_spent' to decide what
54 to explore; they are bookkeeping only. Derive any decision statistic from
55 'question.observed()' instead.
56
57 ## Required branch trajectory and failure interpretation
58
59 For each opened branch, reconstruct its ordered prefix trajectory, not only its
60 latest observation or best score: successful anchor, score trend, regressions,
61 failure/repair sequence, and explored versus remaining depth.
62
63 Before closing or deprioritizing a failed frontier, classify it as
64 hard-unrecoverable, repairable implementation failure, weak-but-underexplored, or
65 repeatedly unpromising after sufficient valid evidence. Output/correctness mismatch,
66 shared-memory/resource limits, and variable/code, mask/layout/shape errors are
67 normally repairable. Do not infer algorithmic failure from one such error.
68 'n_valid == 0' and 'branch_failed_hard(obs)' are signals, not unconditional
69 closure: use 'fail_class' and 'error' to distinguish a repairable zero-valid
70 failure from an environment/dependency failure. 'compile_other' alone is not
71 permanently hard. Classify the current failure episode: a later successful result
72 reopens the branch and cancels closure based only on an earlier failure.
73
74 ## Required batch decision loop
75
76 At each decision round:
77
78 1. Read the prefix, reconstruct trajectories, and close only branches with
79    cumulative evidence of being hard-unrecoverable or repeatedly unpromising.
80 2. Rank legal roots and legal branch frontiers using only prefix-derived signals:
81    successful anchor, parent→child gain, complete branch trajectory, actual success
82    versus failure evidence,
83    failure recoverability, prior repair outcomes, remaining depth, and cross-branch
84    comparison.
85 3. Rank actual repairable failures and underexplored frontiers in deterministic
86    queues using trajectory, recoverability, remaining depth, repeated failures, and
87    beta. A repairable failure retains eligibility unless cumulative evidence lowers
88    its relative priority.
89 4. Build one **dynamic portfolio** batch of independent candidates, up to
90    'question.max_parallelism': exploitation (strong normal refinements),
91    exploration (new roots or underexplored branches), and at most one recovery
92    (an actual repairable failure). When multiple roles are eligible, give
93    exploration and justified recovery representation before filling remaining slots

```

```

94 by priority; adapt this to prefix evidence rather than fixed quotas. Recovery
95 must not displace normal successful refinements or leave workers idle. Never
96 sample randomly, and do not default to a singleton merely because its top
97 candidate is clear.
98 5. Stop only after considering the whole revealed portfolio: active, underexplored,
99 recoverable, unopened, and remaining legal candidates. Do not stop while an
100 eligible high-priority recovery or underexplored candidate remains; every
101 remaining action needs an evidence-based decision to continue, reserve, or close.
102
103 A batch must contain distinct cells that are all legal *before* the call. It may
104 contain several roots and/or one frontier from each opened branch. It must never
105 contain a parent and its child together. Do not use a fixed widen-all / deepen-all
106 wave schedule: adapt batch composition after every revealed prefix.
107
108 Minimal structure:
109
110 from see.policy.api import (
111     LLMDesignedMethod, SimResult, _budget_done, _record_curve, finalize_result,
112 )
113
114 def solve(self, question, budget=None):
115     question.reset()
116     res, closed = SimResult(), set()
117     while not _budget_done(question, budget):
118         prefix = question.observed()
119         update_closed(closed, prefix, question)
120         batch = select_batch(prefix, question, closed)
121         if not batch:
122             break
123         question.probe_batch(
124             batch,
125             on_reveal=lambda _: _record_curve(res, question),
126         )
127     return finalize_result(question, res)
128
129 ## Hard constraints
130
131 - Keep ‘NAME = "OptimalPolicy"’ and implement
132   ‘class OptimalPolicy(LLMDesignedMethod)’ in ‘{method_file}’ only.
133 - **Prefix-only:** decisions may use revealed observations, ‘baseline_score’, legal
134   sets, structural ‘meta’, and helper signals. Never use unrevealed scores, a true
135   optimum, hardcoded winning cell ids, absolute score targets, or internal trace data.
136 - Every prune, widen, deepen, batch, and stop decision must be explainable from the
137   current prefix. Shallow weak scores are not enough to discard a branch: deeper
138   attempts can recover. A repairable latest failure must not erase its historical
139   successful anchor or by itself cause permanent starvation.
140 - Replay calls with ‘budget=None’. Always terminate when no batch is selected; do
141   not assume a budget cap exists.
142 - A selected batch must be legal, have no duplicate ids, and contain at most
143   ‘question.max_parallelism’ cells.
144
145 ## Beta: fixed per run, adaptive across cycles
146
147 Read exactly one scalar in ‘__init__’:
148
149     beta = float(self.config.get("beta", <sensible_default>))
150
151 Beta has three distinct roles. Do not conflate them:
152
153 1. **Within one replay or live episode:** beta is fixed. Route every behavioral
154   threshold through one ‘_schedule(beta) -> dict’. High beta means more width,
155   deeper patience, and weaker pruning. Low beta means fewer probes, earlier
156   stagnation stops, and stronger pruning. Never change beta from observations inside
157   ‘solve()’. Route recovery eligibility, reserve threshold, and waiting through
158   the same schedule: high beta is more patient; low beta remains selective without
159   treating one repairable failure as automatic closure.
160 2. **During offline evaluation:** eval sweeps a fixed beta grid. This measures whether
161   the policy exposes a real attainment/work/parallelism trade-off; it is not online
162   beta adaptation.
163 3. **When proposing the next policy version:** choose the baked-in default beta once,

```

```

164 using evidence from earlier *live* cycles and their beta sweeps. That default will
165 remain fixed throughout the next live exploration episode.
166
167 Keep all thresholds relative to the prefix; never use absolute score cutoffs.
168
169 Use the following cross-cycle default-beta rule. Read the most recent 2-3
170 **live** 'trace_pool/iter*/live_cycle_manifest.json' sidecars (and '_current'
171 when present) for each iteration's final best score and actual baked-in beta. Read
172 the matching archived 'beta_sweep.json' values ('pareto.reward', AUC, parallel
173 penalty, and the per-beta frontier). Scores alone do not establish that beta caused a
174 change, so always use both sources:
175
176 - live best is still improving: keep the prior default beta unless its sweep clearly
177 shows a better nearby beta;
178 - live best has plateaued, and higher beta reaches higher attainment for a reasonable
179 work/parallelism cost in the sweep: raise the default by a small step (about
180 0.1-0.2, clamped to [0, 1]);
181 - a high default beta has already been tried through a plateau, and high-beta sweep
182 points add work without higher attainment: lower it by a small step;
183 - history is insufficient or evidence conflicts: use a moderately exploratory default
184 (about 0.6), rather than pretending the replay ceiling is a live stopping signal.
185
186 The beta sweep is non-degenerate only if beta changes the attainment/work trade-off.
187 It also reveals whether the policy batches. Do not select the default simply as the
188 smallest beta that reaches a frozen trace's known ceiling.
189
190 ## Required next-cycle grid planning
191
192 Every proposed policy **must** implement this deterministic method:
193
194 from see.policy.api import GridPlan, GridPlanningContext
195
196 def plan_grid(self, context: GridPlanningContext) -> GridPlan:
197     ...
198
199 This method runs **before** a new live grid is created. It does not make a
200 within-episode decision and must never inspect a current episode's outcomes.
201 It must always return a non-None 'GridPlan': do not inherit the template
202 stub and do not delegate grid choice to the runner's fallback. When history is
203 empty or insufficient, still return an explicit conservative bootstrap plan
204 derived from the context's fallback/hard-cap fields, with a factual reason.
205
206 'GridPlan(branch_count=W, refine_count=R)' accepts arbitrary integers, not a
207 fixed set of presets. It creates branches '0..W-1' and attempts '0..R'; 'R' is
208 the number of refinements allowed after each root. The runner validates
209 '1 <= W <= context.hard_max_branch_count' and
210 '0 <= R <= context.hard_max_refine_count'. In replay, a requested plan beyond the
211 frozen trace's 'context.trace_branch_count' or 'context.trace_refine_count' is
212 out of support and cannot earn replay reward.
213
214 Use only the prefix-safe facts in 'context':
215
216 - 'history': completed earlier live manifests, including prior planned/effective
217 grids, actual opened width/depth, probe work, decision rounds, scores, and beta;
218 - fallback/hard caps and worker cap;
219 - replay structural support fields. Do not read raw trace outcomes or a current
220 cycle result inside 'plan_grid'.
221
222 Choose width versus depth from evidence, not a default preference:
223
224 - many semantically distinct roots improve early while deeper refinements stall:
225 increase width and reduce/hold depth;
226 - high gains arrive late on a small, repeatable set of directions: reduce/hold width
227 and increase depth;
228 - all explored directions plateau after sufficient depth while meaningful direction
229 classes remain uncovered: increase width;
230 - repeated hard, unrecoverable failures or strongly redundant directions: reduce
231 width and depth conservatively;
232 - conflicting or insufficient history: return an explicit conservative bootstrap
233 plan derived from the context, and state that evidence is insufficient.

```

```

234
235 Include a short, factual ‘‘reason’’ in every plan. ‘‘plan_grid’’ answers
236 how many directions to make available; the direction provider assigns those new
237 roots their directions, and ‘‘solve’’ still decides which legal roots/frontiers to
238 open, refine, prune, or stop. Do not choose roots merely because their branch id is
239 small. The runtime grid is the hard bound: controller thresholds may use less, but
240 can never create branches or attempts beyond the effective plan. Before finishing,
241 verify that the edited ‘‘method.py’’ contains an override of ‘‘plan_grid’’ that
242 returns ‘‘GridPlan(branch_count=..., refine_count=..., reason=...)’’ on every path.
243
244 ## Learn from history without leaking outcomes
245
246 Earlier rounds are in ‘‘{history_dir}/r####.*/’’. Read their policy code and
247 ‘‘proposal_results/beta_sweep.json’’. Start from a strong recent policy, retain
248 mechanisms that raised ‘‘pareto.reward’’, and make a concrete change when progress
249 stalls. A legacy AUC-only sweep is useful code history but is not numerically
250 comparable to the current reward. The baseline under ‘‘{history_dir}/baseline/’’ is
251 a parallel-refine floor to beat.
252
253 Each current-objective round also archives
254 ‘‘proposal_results/policy_execution_traces.jsonl’’: one replay episode per
255 ‘‘(frozen trace, beta)’’. Use it to diagnose general behavior -- serial batches,
256 premature stops, over-pruning, or wasted probes -- from the prefix state, selected
257 batch, and revealed outcomes at each decision round. It is **between-round feedback
258 only**: never read it inside ‘‘solve()’’, and never copy a trace-specific branch,
259 cell id, score, or target into policy logic.
260
261 ‘‘{trace_pool}’’, if present, may be read only outside ‘‘solve()’’. Prefer the
262 ‘‘live_cycle_manifest.json’’ sidecars over raw replay outcomes for the per-iteration
263 live trend. Never copy trace scores, targets, or cell ids into policy logic.
264
265 ## Deliverable
266
267 Write a complete adaptive policy in ‘‘{method_file}’’. Include a short module
268 docstring describing its prefix signals, batch rule, beta schedule, default-beta
269 rationale, grid-planning rule (if implemented), and safeguards against
270 over-pruning, over-stopping, permanent starvation after repairable failures, and
271 serial probes. Before finishing, verify trajectory-based ranking, the stated
272 success semantics, non-automatic zero-valid closure, deterministic recovery
273 competition, and portfolio-level stop.

```

C. Discovered Programs

We provide the complete implementation of the Lasso-path solver discovered by DREAM-RSI. As discussed in Section 4.1, the solver combines strong-rule screening with adaptive Cauchy–Schwarz KKT pruning, disjoint active-set bookkeeping, lazy Gram-matrix construction, and hardware-aware optimizations.

Listing 3 | Complete Lasso-path solver discovered by DREAM-RSI.

```

1 # EVOLVE-BLOCK-START
2
3 CPP_CODE = r'''
4 #define EIGEN_NO_DEBUG
5 #define EIGEN_MPL2_ONLY
6 #define EIGEN_UNROLL_LOOPS
7
8 #include <Eigen/Dense>
9 #include <vector>
10 #include <cstdio>
11 #include <cmath>
12 #include <algorithm>
13 #include <numeric>
14 #include <omp.h>
15 #include <cstdlib>
16 #include <cstring>

```

```

17
18 using Eigen::MatrixXd;
19 using Eigen::VectorXd;
20
21 #if defined(_MSC_VER)
22 #define RESTRICT __restrict
23 #elif defined(__GNUC__) || defined(__clang__)
24 #define RESTRICT __restrict__
25 #else
26 #define RESTRICT
27 #endif
28
29 // High-performance alignment assumption
30 #if defined(__GNUC__) || defined(__clang__)
31 #define ASSUME_ALIGNED(ptr, alignment) (double*)__builtin_assume_aligned((ptr), (alignment))
32 #else
33 #define ASSUME_ALIGNED(ptr, alignment) (ptr)
34 #endif
35
36 // High-performance branch-free soft-thresholding using std::abs and std::copysign
37 static inline double soft_thresh(double z, double gamma) {
38     double abs_z = std::abs(z);
39     double val = abs_z - gamma;
40     return std::copysign(val > 0.0 ? val : 0.0, z);
41 }
42
43 // =====
44 // DISJOINT-PARTITION ACTIVE-SET LASSO PATH SOLVER WITH ALIGNED COLUMN PADDING
45 // =====
46 void solve_active_set(
47     const double* RESTRICT X_padded,
48     int n_padded,
49     int n,
50     int p,
51     const VectorXd& y,
52     const VectorXd& lam_path,
53     const VectorXd& xv,
54     const VectorXd& grad_init,
55     MatrixXd& coef_path, // (p, n_lam) output, pre-zeroed
56     double thresh,      // convergence threshold
57     int maxit)           // max inner loop iterations
58 {
59     const double fn = static_cast<double>(n);
60     const double inv_fn = 1.0 / fn;
61     const double tol = thresh;
62     const int nlam = lam_path.size();
63
64     // Workload-Aware flag for activating Cauchy-Schwarz KKT Pruning
65     const bool use_cs = (p >= 500 && n >= 150);
66
67     // Initial capacity for active set structures - optimized to completely avoid reallocations on
68     // almost all problems
69     int current_capacity = ((std::max(128, std::min(512, p)) + 7) / 8) * 8;
70
71     // Declare raw pointers for 64-byte aligned structures
72     double* G_data = nullptr;
73     double* c_data = nullptr;
74     double* beta_active_data = nullptr;
75     double* xv_active_data = nullptr;
76     double* inv_xv_active_data = nullptr;
77     double* grad_init_active_data = nullptr;
78     double* beta_old_at_start = nullptr;
79
80     double* y_padded = nullptr;
81     double* r_padded = nullptr;
82     double* r_ref_padded = nullptr;
83
84     bool oom = false;
85
86     // Allocate 64-byte aligned arrays

```

```

86   if (posix_memalign((void**)&G_data, 64, static_cast<size_t>(current_capacity) *
87       current_capacity * sizeof(double)) != 0) goto cleanup;
87   if (posix_memalign((void**)&c_data, 64, static_cast<size_t>(current_capacity) * sizeof(double)
88       ) != 0) goto cleanup;
88   if (posix_memalign((void**)&beta_active_data, 64, static_cast<size_t>(current_capacity) *
89       sizeof(double)) != 0) goto cleanup;
89   if (posix_memalign((void**)&xv_active_data, 64, static_cast<size_t>(current_capacity) * sizeof
90       (double)) != 0) goto cleanup;
90   if (posix_memalign((void**)&inv_xv_active_data, 64, static_cast<size_t>(current_capacity) *
91       sizeof(double)) != 0) goto cleanup;
91   if (posix_memalign((void**)&grad_init_active_data, 64, static_cast<size_t>(current_capacity) *
92       sizeof(double)) != 0) goto cleanup;
92   if (posix_memalign((void**)&beta_old_at_start, 64, static_cast<size_t>(current_capacity) *
93       sizeof(double)) != 0) goto cleanup;
93
94   if (posix_memalign((void**)&y_padded, 64, static_cast<size_t>(n_padded) * sizeof(double)) !=
95       0) goto cleanup;
95   if (posix_memalign((void**)&r_padded, 64, static_cast<size_t>(n_padded) * sizeof(double)) !=
96       0) goto cleanup;
96   if (posix_memalign((void**)&r_ref_padded, 64, static_cast<size_t>(n_padded) * sizeof(double))
97       != 0) goto cleanup;
97
98   std::fill(G_data, G_data + static_cast<size_t>(current_capacity) * current_capacity, 0.0);
99   std::fill(c_data, c_data + current_capacity, 0.0);
100  std::fill(beta_active_data, beta_active_data + current_capacity, 0.0);
101  std::fill(xv_active_data, xv_active_data + current_capacity, 0.0);
102  std::fill(inv_xv_active_data, inv_xv_active_data + current_capacity, 0.0);
103  std::fill(grad_init_active_data, grad_init_active_data + current_capacity, 0.0);
104  std::fill(beta_old_at_start, beta_old_at_start + current_capacity, 0.0);
105
106  std::memcpy(y_padded, y.data(), n * sizeof(double));
107  for (int i = n; i < n_padded; ++i) y_padded[i] = 0.0;
108
109  std::memcpy(r_padded, y_padded, n_padded * sizeof(double));
110
111  // Consistently initialize r_ref_padded to y_padded (instead of all zeros) to guarantee 100%
112  tight bounds at start
113  std::memcpy(r_ref_padded, y_padded, n_padded * sizeof(double));
114
115  // Run the solver in a nested block to make goto compile-safe
116  {
117      VectorXd beta = VectorXd::Zero(p);
118
119      std::vector<char> screened(p, 0); // 1 if screened, 0 otherwise
120      std::vector<int> active; // indices of active features (beta != 0)
121      std::vector<int> feat_to_idx(p, -1); // maps feature to index in active set
122
123      // Disjoint tracking partition vectors
124      std::vector<int> unscreened_list(p);
125      std::vector<int> screened_list(p);
126      std::vector<int> screened_to_idx(p, -1);
127
128      int unscreened_size = p;
129      int screened_size = 0;
130      for (int j = 0; j < p; ++j) {
131          unscreened_list[j] = j;
132      }
133
134      VectorXd grad = grad_init; // grad can be modified/overwritten
135
136      // Reference state for Cauchy-Schwarz KKT pruning
137      VectorXd grad_ref;
138      std::vector<double> s;
139      int lambdas_since_reset = 0;
140
141      if (use_cs) {
142          grad_ref = grad_init;
143          s.resize(p);
144          for (int j = 0; j < p; ++j) {
145              s[j] = std::sqrt(xv(j) * inv_fn);

```

```

145     }
146 }
147
148 auto add_active = [&](int j) {
149     if (feat_to_idx[j] != -1) return;
150
151     // 0(1) swap-deletion from screened_list to maintain partition disjointness
152     int idx_in_screened = screened_to_idx[j];
153     if (idx_in_screened >= 0) {
154         int last_j = screened_list[screened_size - 1];
155         screened_list[idx_in_screened] = last_j;
156         screened_to_idx[last_j] = idx_in_screened;
157         --screened_size;
158         screened_to_idx[j] = -1;
159     }
160
161     int old_k = static_cast<int>(active.size());
162     feat_to_idx[j] = old_k;
163     active.push_back(j);
164     int new_k = old_k + 1;
165
166     if (new_k > current_capacity) {
167         int new_capacity = current_capacity * 2;
168
169         double* G_data2 = nullptr;
170         double* c_data2 = nullptr;
171         double* beta_active_data2 = nullptr;
172         double* xv_active_data2 = nullptr;
173         double* inv_xv_active_data2 = nullptr;
174         double* grad_init_active_data2 = nullptr;
175         double* beta_old_at_start2 = nullptr;
176
177         if (posix_memalign((void**)&G_data2, 64, static_cast<size_t>(new_capacity) *
new_capacity * sizeof(double)) != 0) { oom = true; return; }
178         if (posix_memalign((void**)&c_data2, 64, static_cast<size_t>(new_capacity) *
sizeof(double)) != 0) { free(G_data2); oom = true; return; }
179         if (posix_memalign((void**)&beta_active_data2, 64, static_cast<size_t>(
new_capacity) * sizeof(double)) != 0) { free(G_data2); free(c_data2); oom = true; return; }
180         if (posix_memalign((void**)&xv_active_data2, 64, static_cast<size_t>(new_capacity)
* sizeof(double)) != 0) { free(G_data2); free(c_data2); free(beta_active_data2); oom = true;
return; }
181         if (posix_memalign((void**)&inv_xv_active_data2, 64, static_cast<size_t>(
new_capacity) * sizeof(double)) != 0) { free(G_data2); free(c_data2); free(beta_active_data2);
free(xv_active_data2); oom = true; return; }
182         if (posix_memalign((void**)&grad_init_active_data2, 64, static_cast<size_t>(
new_capacity) * sizeof(double)) != 0) { free(G_data2); free(c_data2); free(beta_active_data2);
free(xv_active_data2); free(inv_xv_active_data2); oom = true; return; }
183         if (posix_memalign((void**)&beta_old_at_start2, 64, static_cast<size_t>(
new_capacity) * sizeof(double)) != 0) { free(G_data2); free(c_data2); free(beta_active_data2);
free(xv_active_data2); free(inv_xv_active_data2); free(grad_init_active_data2); oom = true;
return; }
184
185         std::fill(G_data2, G_data2 + static_cast<size_t>(new_capacity) * new_capacity,
0.0);
186
187         if (old_k > 0) {
188             int old_k_padded = (old_k + 7) & ~7;
189             for (int col = 0; col < old_k; ++col) {
190                 double* dest_col = G_data2 + col * new_capacity;
191                 const double* src_col = G_data + col * current_capacity;
192                 #pragma omp simd aligned(dest_col, src_col: 64)
193                 for (int row = 0; row < old_k_padded; ++row) {
194                     dest_col[row] = src_col[row];
195                 }
196             }
197
198             #pragma omp simd aligned(c_data2, c_data: 64)
199             for (int i = 0; i < old_k_padded; ++i) c_data2[i] = c_data[i];
200
201             #pragma omp simd aligned(beta_active_data2, beta_active_data: 64)

```

```

202         for (int i = 0; i < old_k_padded; ++i) beta_active_data2[i] = beta_active_data
203         [i];
204         #pragma omp simd aligned(xv_active_data2, xv_active_data: 64)
205         for (int i = 0; i < old_k_padded; ++i) xv_active_data2[i] = xv_active_data[i];
206
207         #pragma omp simd aligned(inv_xv_active_data2, inv_xv_active_data: 64)
208         for (int i = 0; i < old_k_padded; ++i) inv_xv_active_data2[i] =
209         inv_xv_active_data[i];
210
211         #pragma omp simd aligned(grad_init_active_data2, grad_init_active_data: 64)
212         for (int i = 0; i < old_k_padded; ++i) grad_init_active_data2[i] =
213         grad_init_active_data[i];
214
215         #pragma omp simd aligned(beta_old_at_start2, beta_old_at_start: 64)
216         for (int i = 0; i < old_k_padded; ++i) beta_old_at_start2[i] =
217         beta_old_at_start[i];
218     }
219
220     free(G_data);
221     free(c_data);
222     free(beta_active_data);
223     free(xv_active_data);
224     free(inv_xv_active_data);
225     free(grad_init_active_data);
226     free(beta_old_at_start);
227
228     G_data = G_data2;
229     c_data = c_data2;
230     beta_active_data = beta_active_data2;
231     xv_active_data = xv_active_data2;
232     inv_xv_active_data = inv_xv_active_data2;
233     grad_init_active_data = grad_init_active_data2;
234     beta_old_at_start = beta_old_at_start2;
235     current_capacity = new_capacity;
236 }
237
238 // SIMD 4x Register-Blocked Lazy Gram Precomputation (reduces column loads by 75%)
239 const double* RESTRICT col_j = ASSUME_ALIGNED(X_padded + j * n_padded, 64);
240 const bool run_parallel_lazy = (old_k >= 64 && static_cast<size_t>(n_padded) * old_k
241 >= 150000);
242
243 #pragma omp parallel for schedule(static) if(run_parallel_lazy)
244 for (int i = 0; i < (old_k / 4) * 4; i += 4) {
245     const double* RESTRICT col0 = ASSUME_ALIGNED(X_padded + active[i] * n_padded, 64);
246     const double* RESTRICT col1 = ASSUME_ALIGNED(X_padded + active[i+1] * n_padded,
247 64);
248     const double* RESTRICT col2 = ASSUME_ALIGNED(X_padded + active[i+2] * n_padded,
249 64);
250     const double* RESTRICT col3 = ASSUME_ALIGNED(X_padded + active[i+3] * n_padded,
251 64);
252
253     double sum0 = 0.0, sum1 = 0.0, sum2 = 0.0, sum3 = 0.0;
254     #pragma omp simd reduction(+:sum0, sum1, sum2, sum3) aligned(col_j, col0, col1,
255 col2, col3: 64)
256     for (int k = 0; k < n_padded; ++k) {
257         double vj = col_j[k];
258         sum0 += vj * col0[k];
259         sum1 += vj * col1[k];
260         sum2 += vj * col2[k];
261         sum3 += vj * col3[k];
262     }
263
264     double r0 = sum0 * inv_fn;
265     double r1 = sum1 * inv_fn;
266     double r2 = sum2 * inv_fn;
267     double r3 = sum3 * inv_fn;
268
269     G_data[old_k * current_capacity + i] = r0;
270     G_data[i * current_capacity + old_k] = r0;

```

```

263         G_data[old_k * current_capacity + i + 1] = r1;
264         G_data[(i + 1) * current_capacity + old_k] = r1;
265
266         G_data[old_k * current_capacity + i + 2] = r2;
267         G_data[(i + 2) * current_capacity + old_k] = r2;
268
269         G_data[old_k * current_capacity + i + 3] = r3;
270         G_data[(i + 3) * current_capacity + old_k] = r3;
271     }
272
273     for (int i = (old_k / 4) * 4; i < old_k; ++i) {
274         const double* RESTRICT col_act = ASSUME_ALIGNED(X_padded + active[i] * n_padded,
275 64);
276
277         double dot_val = 0.0;
278         #pragma omp simd reduction(+:dot_val) aligned(col_j, col_act: 64)
279         for (int k = 0; k < n_padded; ++k) {
280             dot_val += col_j[k] * col_act[k];
281         }
282         dot_val *= inv_fn;
283         G_data[old_k * current_capacity + i] = dot_val;
284         G_data[i * current_capacity + old_k] = dot_val;
285     }
286     G_data[old_k * current_capacity + old_k] = xv(j); // xv(j) is already scaled by inv_fn
287
288     // Zero-0(n) initial correlation computation
289     double sum_val = 0.0;
290     const double* RESTRICT G_col = ASSUME_ALIGNED(G_data + old_k * current_capacity, 64);
291     const double* RESTRICT beta_act = ASSUME_ALIGNED(beta_active_data, 64);
292     #pragma omp simd reduction(+:sum_val) aligned(G_col, beta_act: 64)
293     for (int i = 0; i < old_k; ++i) {
294         sum_val += G_col[i] * beta_act[i];
295     }
296     c_data[old_k] = grad_init(j) - sum_val;
297
298     xv_active_data[old_k] = xv(j);
299     inv_xv_active_data[old_k] = 1.0 / xv(j);
300     grad_init_active_data[old_k] = grad_init[j];
301     beta_active_data[old_k] = 0.0;
302 };
303
304 double prev_lam = 0.0;
305
306 // Preallocate vectors to avoid repeated heap allocation
307 std::vector<int> to_activate;
308 std::vector<int> screened_violators;
309 std::vector<int> unscreened_violators;
310 std::vector<int> to_compute;
311
312 to_activate.reserve(p);
313 screened_violators.reserve(p);
314 unscreened_violators.reserve(p);
315 if (use_cs) {
316     to_compute.reserve(p);
317 }
318
319 for (int li = 0; li < nlam; ++li) {
320     const double lam = lam_path(li);
321     const double tlam = 2.0 * lam - prev_lam;
322
323     // ---- Step 1: Strong-rule screening (with 0(1) swap-deletion) ----
324     double* RESTRICT grad_ptr = grad.data();
325     for (int i = 0; i < unscreened_size; ) {
326         int j = unscreened_list[i];
327         if (std::abs(grad_ptr[j]) > tlam) {
328             screened[j] = 1;
329             screened_to_idx[j] = screened_size;
330             screened_list[screened_size++] = j;
331             unscreened_list[i] = unscreened_list[--unscreened_size];
332         } else {

```

```

332         ++i;
333     }
334 }
335
336 // ---- Step 2: Outer loop ----
337 int nlp = 0;
338 while (true) {
339     // 2a. Identify violating features among screened features
340     to_activate.clear();
341     const double KKT_bound_screen = lam * (1.0 + 1e-9);
342     for (int i = 0; i < screened_size; ++i) {
343         int j = screened_list[i];
344         // At this point, screened_list only contains non-active screened features.
345         // Absolutely no feat_to_idx branches needed!
346         if (std::abs(grad_ptr[j]) > KKT_bound_screen) {
347             to_activate.push_back(j);
348         }
349     }
350
351     // If some screened features violate KKT, add them to active set
352     if (!to_activate.empty()) {
353         for (int j : to_activate) {
354             add_active(j);
355             if (oom) goto cleanup;
356         }
357     }
358
359     // 2b. CD over active set until convergence
360     int active_size = static_cast<int>(active.size());
361
362     // Save beta at the start of the outer iteration to track changes
363     if (active_size > 0) {
364         int active_size_padded = (active_size + 7) & ~7;
365         #pragma omp simd aligned(beta_old_at_start, beta_active_data: 64)
366         for (int i = 0; i < active_size_padded; ++i) {
367             beta_old_at_start[i] = beta_active_data[i];
368         }
369     }
370
371     if (active_size > 0) {
372         double dmax = tol; // Ensure at least one sweep
373         while (dmax >= tol && nlp < maxit) {
374             ++nlp;
375             dmax = 0.0;
376             for (int idx = 0; idx < active_size; ++idx) {
377                 const double bj_old = beta_active_data[idx];
378                 // Division-free gradient calculation
379                 const double gj = c_data[idx] + bj_old * xv_active_data[idx];
380                 const double bj_new = soft_thresh(gj, lam) * inv_xv_active_data[idx];
381                 if (bj_new == bj_old) continue;
382                 const double delta = bj_new - bj_old;
383                 beta_active_data[idx] = bj_new;
384
385                 // Extremely fast SIMD cache update (padded up to a multiple of 8)
386                 int active_size_padded = (active_size + 7) & ~7;
387                 double* RESTRICT c_ptr = ASSUME_ALIGNED(c_data, 64);
388                 const double* RESTRICT G_col_ptr = ASSUME_ALIGNED(G_data + idx *
current_capacity, 64);
389                 #pragma omp simd aligned(c_ptr, G_col_ptr: 64)
390                 for (int i = 0; i < active_size_padded; ++i) {
391                     c_ptr[i] -= delta * G_col_ptr[i];
392                 }
393
394                 const double ch = xv_active_data[idx] * delta * delta;
395                 if (ch > dmax) dmax = ch;
396             }
397         }
398     }
399
400     // Safety limit check

```

```

401         if (nlp >= maxit) break;
402
403         // Incremental O(n) residual update & any_changed check (Raw-Pointer hand-
vectorized loop)
404         bool any_changed = false;
405         if (active_size > 0) {
406             double* RESTRICT r_ptr = ASSUME_ALIGNED(r_padded, 64);
407             for (int idx = 0; idx < active_size; ++idx) {
408                 const double delta = beta_active_data[idx] - beta_old_at_start[idx];
409                 if (delta != 0.0) {
410                     const double* RESTRICT col_ptr = ASSUME_ALIGNED(X_padded + active[idx]
* n_padded, 64);
411                     #pragma omp simd aligned(r_ptr, col_ptr: 64)
412                     for (int i = 0; i < n_padded; ++i) {
413                         r_ptr[i] -= delta * col_ptr[i];
414                     }
415                     any_changed = true;
416                 }
417             }
418         }
419
420         // O(k^2) exact re-sync of correlation cache c (Sparse-Skipping Custom Loop)
421         if (any_changed && active_size > 0) {
422             int active_size_padded = (active_size + 7) & ~7;
423             #pragma omp simd aligned(c_data, grad_init_active_data: 64)
424             for (int i = 0; i < active_size_padded; ++i) {
425                 c_data[i] = grad_init_active_data[i];
426             }
427             for (int j = 0; j < active_size; ++j) {
428                 const double bj = beta_active_data[j];
429                 if (bj != 0.0) {
430                     const double* RESTRICT G_col = ASSUME_ALIGNED(G_data + j *
current_capacity, 64);
431                     double* RESTRICT c_ptr = ASSUME_ALIGNED(c_data, 64);
432                     #pragma omp simd aligned(c_ptr, G_col: 64)
433                     for (int i = 0; i < active_size_padded; ++i) {
434                         c_ptr[i] -= bj * G_col[i];
435                     }
436                 }
437             }
438         }
439
440         // 2c. Robust Two-Stage KKT check
441         bool screened_kkt_ok = true;
442         screened_violators.clear();
443         const double KKT_bound = lam * (1.0 + 1e-9);
444
445         // SIMD 4x Register-Blocked Screened KKT Checks (reduces residual vector loads by
75%)
446         const double* RESTRICT r_ptr = ASSUME_ALIGNED(r_padded, 64);
447         const bool run_parallel_screened = (static_cast<size_t>(n_padded) * screened_size
>= 150000);
448
449         #pragma omp parallel for schedule(static) if(run_parallel_screened)
450         for (int i = 0; i < (screened_size / 4) * 4; i += 4) {
451             int j0 = screened_list[i];
452             int j1 = screened_list[i+1];
453             int j2 = screened_list[i+2];
454             int j3 = screened_list[i+3];
455
456             const double* RESTRICT col0 = ASSUME_ALIGNED(X_padded + j0 * n_padded, 64);
457             const double* RESTRICT col1 = ASSUME_ALIGNED(X_padded + j1 * n_padded, 64);
458             const double* RESTRICT col2 = ASSUME_ALIGNED(X_padded + j2 * n_padded, 64);
459             const double* RESTRICT col3 = ASSUME_ALIGNED(X_padded + j3 * n_padded, 64);
460
461             double sum0 = 0.0, sum1 = 0.0, sum2 = 0.0, sum3 = 0.0;
462             #pragma omp simd reduction(+:sum0, sum1, sum2, sum3) aligned(r_ptr, col0, col1
, col2, col3: 64)
463             for (int k = 0; k < n_padded; ++k) {
464                 double rk = r_ptr[k];

```

```

465         sum0 += rk * col0[k];
466         sum1 += rk * col1[k];
467         sum2 += rk * col2[k];
468         sum3 += rk * col3[k];
469     }
470     grad_ptr[j0] = sum0 * inv_fn;
471     grad_ptr[j1] = sum1 * inv_fn;
472     grad_ptr[j2] = sum2 * inv_fn;
473     grad_ptr[j3] = sum3 * inv_fn;
474 }
475
476 for (int i = (screened_size / 4) * 4; i < screened_size; ++i) {
477     int j = screened_list[i];
478     const double* RESTRICT col_ptr = ASSUME_ALIGNED(X_padded + j * n_padded, 64);
479     double dot_val = 0.0;
480     #pragma omp simd reduction(+:dot_val) aligned(r_ptr, col_ptr: 64)
481     for (int k = 0; k < n_padded; ++k) {
482         dot_val += col_ptr[k] * r_ptr[k];
483     }
484     grad_ptr[j] = dot_val * inv_fn;
485 }
486
487 for (int i = 0; i < screened_size; ++i) {
488     int j = screened_list[i];
489     if (std::abs(grad_ptr[j]) > KKT_bound) {
490         screened_violators.push_back(j);
491         screened_kkt_ok = false;
492     }
493 }
494
495 if (!screened_kkt_ok) {
496     // Add screened violators to active set and run CD again
497     for (int j : screened_violators) {
498         add_active(j);
499         if (oom) goto cleanup;
500     }
501     continue; // Skip full KKT check, go back to CD
502 }
503
504 // Only perform full KKT check on unscreened features if screened is 100% OK
505 bool full_kkt_ok = true;
506 unscreened_violators.clear();
507
508 if (use_cs) {
509     // Dual-Phase Adaptive Cauchy-Schwarz KKT Pruning!
510     double d2 = 0.0;
511     const double* RESTRICT r_curr_ptr = ASSUME_ALIGNED(r_padded, 64);
512     const double* RESTRICT r_ref_ptr = ASSUME_ALIGNED(r_ref_padded, 64);
513     #pragma omp simd reduction(+:d2) aligned(r_curr_ptr, r_ref_ptr: 64)
514     for (int k = 0; k < n_padded; ++k) {
515         double diff = r_curr_ptr[k] - r_ref_ptr[k];
516         d2 += diff * diff;
517     }
518     double d = std::sqrt(d2);
519
520     const double* RESTRICT grad_ref_ptr = grad_ref.data();
521     const double* RESTRICT s_ptr = s.data();
522     const int* RESTRICT unscreened_ptr = unscreened_list.data();
523
524     to_compute.clear();
525     for (int i = 0; i < unscreened_size; ++i) {
526         int j = unscreened_ptr[i];
527         double bound = std::abs(grad_ref_ptr[j]) + s_ptr[j] * d;
528         if (bound > KKT_bound) {
529             to_compute.push_back(j);
530         }
531     }
532
533     int num_to_compute = to_compute.size();
534     bool did_reset = false;

```

```

535         if (num_to_compute > 0.3 * p || lambdas_since_reset >= 8) {
536             // Drift is too large or reset interval reached, do a full reset (SIMD 4x
537             Register-Blocked)
538             const bool run_parallel_reset = (static_cast<size_t>(n_padded) *
unscreened_size >= 150000);
539             #pragma omp parallel for schedule(static) if(run_parallel_reset)
540             for (int i = 0; i < (unscreened_size / 4) * 4; i += 4) {
541                 int j0 = unscreened_list[i];
542                 int j1 = unscreened_list[i+1];
543                 int j2 = unscreened_list[i+2];
544                 int j3 = unscreened_list[i+3];
545
546                 const double* RESTRICT col0 = ASSUME_ALIGNED(X_padded + j0 * n_padded,
64);
547                 const double* RESTRICT col1 = ASSUME_ALIGNED(X_padded + j1 * n_padded,
64);
548                 const double* RESTRICT col2 = ASSUME_ALIGNED(X_padded + j2 * n_padded,
64);
549                 const double* RESTRICT col3 = ASSUME_ALIGNED(X_padded + j3 * n_padded,
64);
550                 const double* RESTRICT r_ptr_exact = ASSUME_ALIGNED(r_padded, 64);
551
552                 double sum0 = 0.0, sum1 = 0.0, sum2 = 0.0, sum3 = 0.0;
553                 #pragma omp simd reduction(+:sum0, sum1, sum2, sum3) aligned(
r_ptr_exact, col0, col1, col2, col3: 64)
554                 for (int k = 0; k < n_padded; ++k) {
555                     double rk = r_ptr_exact[k];
556                     sum0 += rk * col0[k];
557                     sum1 += rk * col1[k];
558                     sum2 += rk * col2[k];
559                     sum3 += rk * col3[k];
560                 }
561                 grad_ptr[j0] = sum0 * inv_fn;
562                 grad_ptr[j1] = sum1 * inv_fn;
563                 grad_ptr[j2] = sum2 * inv_fn;
564                 grad_ptr[j3] = sum3 * inv_fn;
565             }
566
567             for (int i = (unscreened_size / 4) * 4; i < unscreened_size; ++i) {
568                 int j = unscreened_list[i];
569                 const double* RESTRICT col_ptr = ASSUME_ALIGNED(X_padded + j *
n_padded, 64);
570
571                 const double* RESTRICT r_ptr_exact = ASSUME_ALIGNED(r_padded, 64);
572                 double sum = 0.0;
573                 #pragma omp simd reduction(+:sum) aligned(r_ptr_exact, col_ptr: 64)
574                 for (int k = 0; k < n_padded; ++k) {
575                     sum += r_ptr_exact[k] * col_ptr[k];
576                 }
577                 grad_ptr[j] = sum * inv_fn;
578             }
579
580             std::memcpy(r_ref_padded, r_padded, n_padded * sizeof(double));
581
582             double* RESTRICT grad_ref_ptr_writable = grad_ref.data();
583             #pragma omp parallel for schedule(static) if(unscreened_size >= 2048)
584             for (int i = 0; i < unscreened_size; ++i) {
585                 int j = unscreened_ptr[i];
586                 grad_ref_ptr_writable[j] = grad_ptr[j];
587             }
588             lambdas_since_reset = 0;
589             did_reset = true;
590         } else {
591             // Compute exact gradients only for the tiny unpruned subset (SIMD 4x
Register-Blocked)
592             const bool run_parallel_comp = (num_to_compute >= 32 && static_cast<size_t>
>(n_padded) * num_to_compute >= 150000);
593             #pragma omp parallel for schedule(static) if(run_parallel_comp)
594             for (int k = 0; k < (num_to_compute / 4) * 4; k += 4) {
595                 int j0 = to_compute[k];

```

```

595         int j1 = to_compute[k+1];
596         int j2 = to_compute[k+2];
597         int j3 = to_compute[k+3];
598
599         const double* RESTRICT col0 = ASSUME_ALIGNED(X_padded + j0 * n_padded,
600         64);
601         const double* RESTRICT col1 = ASSUME_ALIGNED(X_padded + j1 * n_padded,
602         64);
603         const double* RESTRICT col2 = ASSUME_ALIGNED(X_padded + j2 * n_padded,
604         64);
605         const double* RESTRICT col3 = ASSUME_ALIGNED(X_padded + j3 * n_padded,
606         64);
607         const double* RESTRICT r_ptr_exact = ASSUME_ALIGNED(r_padded, 64);
608
609         double sum0 = 0.0, sum1 = 0.0, sum2 = 0.0, sum3 = 0.0;
610         #pragma omp simd reduction(+:sum0, sum1, sum2, sum3) aligned(
611         r_ptr_exact, col0, col1, col2, col3: 64)
612         for (int m = 0; m < n_padded; ++m) {
613             double rk = r_ptr_exact[m];
614             sum0 += rk * col0[m];
615             sum1 += rk * col1[m];
616             sum2 += rk * col2[m];
617             sum3 += rk * col3[m];
618         }
619         grad_ptr[j0] = sum0 * inv_fn;
620         grad_ptr[j1] = sum1 * inv_fn;
621         grad_ptr[j2] = sum2 * inv_fn;
622         grad_ptr[j3] = sum3 * inv_fn;
623     }
624
625     for (int k = (num_to_compute / 4) * 4; k < num_to_compute; ++k) {
626         int j = to_compute[k];
627         const double* RESTRICT col_ptr = ASSUME_ALIGNED(X_padded + j *
628         n_padded, 64);
629
630         const double* RESTRICT r_ptr_exact = ASSUME_ALIGNED(r_padded, 64);
631         double sum = 0.0;
632         #pragma omp simd reduction(+:sum) aligned(r_ptr_exact, col_ptr: 64)
633         for (int m = 0; m < n_padded; ++m) {
634             sum += r_ptr_exact[m] * col_ptr[m];
635         }
636         grad_ptr[j] = sum * inv_fn;
637     }
638
639     for (int i = 0; i < unscreened_size; ) {
640         int j = unscreened_list[i];
641         if (std::abs(grad_ptr[j]) > KKT_bound) {
642             screened[j] = 1;
643             unscreened_violators.push_back(j);
644             screened_to_idx[j] = screened_size;
645             screened_list[screened_size++] = j;
646             unscreened_list[i] = unscreened_list[--unscreened_size];
647             full_kkt_ok = false;
648         } else {
649             ++i;
650         }
651     }
652
653     if (full_kkt_ok) {
654         if (!did_reset) {
655             lambdas_since_reset++;
656         }
657     }
658 } else {
659     // Standard, clean KKT check without CS pruning overhead on small/medium
660     problems (SIMD 4x Register-Blocked)
661     const bool run_parallel_uns_std = (static_cast<size_t>(n_padded) *
662     unscreened_size >= 150000);
663     #pragma omp parallel for schedule(static) if(run_parallel_uns_std)
664     for (int i = 0; i < (unscreened_size / 4) * 4; i += 4) {

```

```

657         int j0 = unscreened_list[i];
658         int j1 = unscreened_list[i+1];
659         int j2 = unscreened_list[i+2];
660         int j3 = unscreened_list[i+3];
661
662         const double* RESTRICT col0 = ASSUME_ALIGNED(X_padded + j0 * n_padded, 64)
;
663         const double* RESTRICT col1 = ASSUME_ALIGNED(X_padded + j1 * n_padded, 64)
;
664         const double* RESTRICT col2 = ASSUME_ALIGNED(X_padded + j2 * n_padded, 64)
;
665         const double* RESTRICT col3 = ASSUME_ALIGNED(X_padded + j3 * n_padded, 64)
;
666         const double* RESTRICT r_ptr = ASSUME_ALIGNED(r_padded, 64);
667
668         double sum0 = 0.0, sum1 = 0.0, sum2 = 0.0, sum3 = 0.0;
669         #pragma omp simd reduction(+:sum0, sum1, sum2, sum3) aligned(r_ptr, col0,
col1, col2, col3: 64)
670         for (int k = 0; k < n_padded; ++k) {
671             double rk = r_ptr[k];
672             sum0 += rk * col0[k];
673             sum1 += rk * col1[k];
674             sum2 += rk * col2[k];
675             sum3 += rk * col3[k];
676         }
677         grad_ptr[j0] = sum0 * inv_fn;
678         grad_ptr[j1] = sum1 * inv_fn;
679         grad_ptr[j2] = sum2 * inv_fn;
680         grad_ptr[j3] = sum3 * inv_fn;
681     }
682
683     for (int i = (unscreened_size / 4) * 4; i < unscreened_size; ++i) {
684         int j = unscreened_list[i];
685         const double* RESTRICT col_ptr = ASSUME_ALIGNED(X_padded + j * n_padded,
64);
686
687         const double* RESTRICT r_ptr = ASSUME_ALIGNED(r_padded, 64);
688         double sum = 0.0;
689         #pragma omp simd reduction(+:sum) aligned(r_ptr, col_ptr: 64)
690         for (int k = 0; k < n_padded; ++k) {
691             sum += r_ptr[k] * col_ptr[k];
692         }
693         grad_ptr[j] = sum * inv_fn;
694     }
695
696     for (int i = 0; i < unscreened_size; ) {
697         int j = unscreened_list[i];
698         if (std::abs(grad_ptr[j]) > KKT_bound) {
699             screened[j] = 1;
700             unscreened_violators.push_back(j);
701             screened_to_idx[j] = screened_size;
702             screened_list[screened_size++] = j;
703             unscreened_list[i] = unscreened_list[--unscreened_size];
704             full_kkt_ok = false;
705         } else {
706             ++i;
707         }
708     }
709
710     if (full_kkt_ok) {
711         break; // Converged completely!
712     }
713
714     // Add unscreened violators to active set
715     for (int j : unscreened_violators) {
716         add_active(j);
717         if (oom) goto cleanup;
718     }
719 }
720

```

```

721         // Synchronize beta with beta_active and save coefficients
722         for (size_t idx = 0; idx < active.size(); ++idx) {
723             beta[active[idx]] = beta_active_data[idx];
724         }
725         coef_path.col(1i) = beta;
726         prev_lam = lam;
727     }
728 }
729
730 cleanup:
731     if (G_data) free(G_data);
732     if (c_data) free(c_data);
733     if (beta_active_data) free(beta_active_data);
734     if (xv_active_data) free(xv_active_data);
735     if (inv_xv_active_data) free(inv_xv_active_data);
736     if (grad_init_active_data) free(grad_init_active_data);
737     if (beta_old_at_start) free(beta_old_at_start);
738     if (y_padded) free(y_padded);
739     if (r_padded) free(r_padded);
740     if (r_ref_padded) free(r_ref_padded);
741 }
742
743 int main() {
744     int32_t n, p, n_lambda;
745     if (fread(&n, sizeof(int32_t), 1, stdin) != 1) return 1;
746     if (fread(&p, sizeof(int32_t), 1, stdin) != 1) return 1;
747     if (fread(&n_lambda, sizeof(int32_t), 1, stdin) != 1) return 1;
748
749     // X arrives row-major. Allocate RowMajor matrix to read the bytes directly!
750     Eigen::Matrix<double, Eigen::Dynamic, Eigen::Dynamic, Eigen::RowMajor> X_row(n, p);
751     if (fread(X_row.data(), sizeof(double), static_cast<size_t>(n) * p, stdin)
752         != static_cast<size_t>(n) * p) return 1;
753
754     // Pad row dimension of X to the multiple of 8 (guarantees perfect alignment for each column)
755     int n_padded = ((n + 7) / 8) * 8;
756     double* X_padded = nullptr;
757     if (posix_memalign((void**)&X_padded, 64, static_cast<size_t>(n_padded) * p * sizeof(double))
758         != 0) return 1;
759
760     VectorXd y(n);
761     if (fread(y.data(), sizeof(double), n, stdin) != static_cast<size_t>(n)) return 1;
762
763     VectorXd lam_path(n_lambda);
764     if (fread(lam_path.data(), sizeof(double), n_lambda, stdin)
765         != static_cast<size_t>(n_lambda)) return 1;
766
767     MatrixXd coef_path = MatrixXd::Zero(p, n_lambda);
768
769     VectorXd xv(p);
770     VectorXd grad_init(p);
771
772     const double* RESTRICT y_ptr = y.data();
773     const double inv_fn = 1.0 / n;
774
775     // 2D Cache-Blocked parallel Fused Transposition-Precomputation-Padding (FTPP)
776     // Avoids separate allocation/std::fill overhead of X_padded and completely saves a full pass
777     // reading X!
778     #pragma omp parallel
779     {
780         int nthreads = omp_get_num_threads();
781         int tid = omp_get_thread_num();
782
783         // Static partition of columns j to completely prevent thread false-sharing
784         int j_per_thread = (p + nthreads - 1) / nthreads;
785         int sj = tid * j_per_thread;
786         int ej = std::min(sj + j_per_thread, p);
787
788         if (sj < ej) {
789             const int col_block = 64;
790             const int row_block = 64;

```

```

789     for (int bj = sj; bj < ej; bj += col_block) {
790         int lim_j = std::min(bj + col_block, ej);
791
792         double local_xx[64] = {0.0};
793         double local_xy[64] = {0.0};
794
795         for (int bi = 0; bi < n; bi += row_block) {
796             int lim_i = std::min(bi + row_block, n);
797             for (int j = bj; j < lim_j; ++j) {
798                 int local_j = j - bj;
799                 double* RESTRICT dest = X_padded + j * n_padded;
800                 const double* RESTRICT src = X_row.data() + j;
801
802                 double sum_xx = 0.0;
803                 double sum_xy = 0.0;
804                 #pragma omp simd reduction(+:sum_xx, sum_xy)
805                 for (int i = bi; i < lim_i; ++i) {
806                     double val = src[i * p];
807                     dest[i] = val;
808                     sum_xx += val * val;
809                     sum_xy += val * y_ptr[i];
810                 }
811                 local_xx[local_j] += sum_xx;
812                 local_xy[local_j] += sum_xy;
813             }
814         }
815
816         // Set the padded elements of each column to 0.0, and store precomputed xv and
grad_init
817         for (int j = bj; j < lim_j; ++j) {
818             double* RESTRICT dest = X_padded + j * n_padded;
819             for (int i = n; i < n_padded; ++i) {
820                 dest[i] = 0.0;
821             }
822             xv(j) = local_xx[j - bj] * inv_fn;
823             grad_init(j) = local_xy[j - bj] * inv_fn;
824         }
825     }
826 }
827
828 // Immediately free memory of X_row to minimize memory footprint
829 X_row.resize(0, 0);
830
831 const double thresh = 1e-9;
832 const int maxit = 100000;
833
834 solve_active_set(X_padded, n_padded, n, p, y, lam_path, xv, grad_init, coef_path, thresh,
835 maxit);
836
837 fwrite(coef_path.data(), sizeof(double),
838        static_cast<size_t>(p) * n_lambda, stdout);
839
840 free(X_padded);
841 return 0;
842 }
843 '''
844
845 COMPILE_FLAGS = ["-fopenmp", "-ffast-math"]
846
847 # EVOLVE-BLOCK-END

```